

Abandoned Object Detection in Video

A BACHELOR'S THESIS

*submitted in partial fulfillment
of the requirements for the award of the degree
of*

BACHELOR OF TECHNOLOGY

in

**INFORMATION TECHNOLOGY
(B.Tech in IT)**

Submitted by

Abhineet Kumar Singh (IIT2009148)

Under the Guidance of:

Dr. Anupam Agrawal

Professor
IIIT-Allahabad



**INDIAN INSTITUTE OF INFORMATION TECHNOLOGY
ALLAHABAD – 211 012 (INDIA)**

July, 2013

CANDIDATE’S DECLARATION

I hereby declare that the work presented in this thesis entitled “**Abandoned object detection in video**”, submitted in the partial fulfillment of the degree of Bachelor of Technology (B.Tech), in Information Technology at Indian Institute of Information Technology, Allahabad, is an authentic record of my original work carried out under the guidance of **Prof. Anupam Agrawal**.

Due acknowledgements have been made in the text of the thesis to all other material used. This thesis work was done in full compliance with the requirements and constraints of the prescribed curriculum.

Place: Allahabad

Date:

Abhineet Kumar Singh

R. No. IIT2009148

CERTIFICATE FROM SUPERVISOR

I do hereby recommend that the thesis work prepared under my supervision by Abhineet Kumar Singh titled “Abandoned object detection in video” be accepted in the partial fulfillment of the requirements of the degree of Bachelor of Technology in Information Technology for Examination.

Date:

Dr. Anupam Agrawal

Place: Allahabad

Professor, IIITA

Committee for Evaluation of the Thesis

ACKNOWLEDGEMENTS

The author would like to express his sincere gratitude to his project supervisor Prof. Anupam Agrawal for his support and guidance that have been crucial to the successful completion of this project.

Place: Allahabad
Date:

Abhineet Kumar Singh
B Tech Final Year, IIITA

ABSTRACT

Abandoned object detection is one of the most practically useful areas of computer vision due to its application in automated video surveillance systems for the detection of suspicious activities that might endanger public safety, especially in crowded places like airports, railway stations, shopping malls, movie theatres and the like. An abandoned object is defined as one that has been lying stationary at a certain place with no apparent human attendance for an extended period of time. Such objects are usually inconspicuous commonplace objects that people often carry around including backpacks, suitcases and boxes. Detection of abandoned objects is of prime importance in uncovering and forestalling terrorist activities since it is a reasonable supposition that an abandoned object, if left behind on purpose, may be hiding dangerous items like explosives.

The present work is an attempt to create a flexible and modular framework that can be used to experiment with several different methods for each stage of the overall task of detecting abandoned and removed objects in a video stream. Several existing methods have been implemented for each of these stages and integrated into the system in a way that makes it possible to switch between them in real time. This enables the user to observe and compare the performance of different methods of solving the same problem and choose the one most suited to any given scenario. The system has also been designed to allow new methods to be added for any stage of the system with minimum programming effort.

Table of Contents

1. Introduction.....	1
1.1 Currently existing technologies.....	1
1.2 Analysis of previous research in this area.....	3
1.3 Problem definition and scope.....	4
1.4 Formulation of the present problem.....	5
1.5 Organization of the thesis.....	6
2. Description of Hardware and Software Used.....	6
2.1 Hardware.....	6
2.2 Software.....	6
3. Theoretical Tools – Analysis and Development.....	7
3.1 Pre-processing.....	8
3.2 Background Modeling and Subtraction.....	11
3.3 Foreground analysis.....	14
3.4 Blob extraction.....	18
3.5 Blob tracking.....	18
3.6 Abandonment analysis.....	20
3.7 Blob filtering.....	22
4. Development of Software	22
4.1 User interface module.....	25
4.2 Pre-processing module.....	25
4.3 BGS module.....	26
4.4 Foreground analysis module.....	27
4.5 Blob extraction module	27
4.6 Blob tracking module.....	27
4.7 Abandonment analysis module.....	29
4.8 Blob filtering module.....	29
5. Testing and Analysis.....	31
5.1 Setup.....	31
5.2 Datasets.....	31
5.3 Result summary.....	34
5.4 Result analysis.....	36
5.5 Comparison with existing systems.....	37
6. Conclusions and Future Work.....	38
Appendix - Explanation of the Source Code.....	39
References.....	48

1. Introduction

An automatic abandoned object detection system typically uses a combination of background subtraction and object tracking to look for certain pre defined patterns of activity that occur when an object is left behind by its owner. Though humans are much better at this task than even state of the art systems, it is often practically unfeasible to employ enough manpower to continuously monitor each one of the very large number of cameras that are required in a large scale surveillance scenario. An automatic detection system therefore helps to complement the available manpower by serving as both a standalone monitoring system for less critical areas and also integrated into manually monitored cameras so that it can detect any drops that the human may have missed.

1.1 Currently existing technologies: Most existing techniques of abandoned (and removed) object detection employ a modular approach with several independent steps where the output of each step serves as the input for the next one. Many efficient algorithms exist for carrying out each of these steps and any single complete AOD system has to address the problem of finding a suitable combination of algorithms to suit a specific scenario. Following is a brief description of these steps and related methods, in the order they are carried out:

1.1.1 Background Modeling and Subtraction (BGS): This stage creates a dynamic model of the scene background and subtracts it from each incoming frame to detect the current foreground regions. The output of this stage is usually a mask depicting pixels in the current frame that do not match the current background model. Some popular background modeling techniques include adaptive medians [1], running averages [2], mixture of Gaussians [3, 4], kernel density estimators [5, 6], Eigen-backgrounds [7] and mean-shift based estimation [8]. There also exist methods that employ dual backgrounds [9] or dual foregrounds [10] for this purpose. The BGS step often utilizes feedback from the object tracking stage to improve its performance.

1.1.2 Foreground Analysis: The BGS step is often unable to adapt to sudden changes in the scene (of lighting, etc.) since the background model is typically updated slowly. It might also confuse parts of a foreground object as background if their appearance happens to be similar to the corresponding background, thus causing a single object to be split into multiple foreground blobs. In addition, certain foreground areas, while being detected correctly, are not of interest for

further processing. The above factors necessitate an additional refinement stage to remove both false foreground regions, caused by factors like background state changes and lighting variations, as well as correct but uninteresting foreground areas like shadows.

Several methods exist for detecting sudden lighting changes, ranging from simple gradient and texture based approaches [11, 12] to those that utilize complex lighting invariant features combined with binary classifiers like support vector machines [13]. Shadow detection is usually carried out by performing a pixel-by-pixel comparison between the current frame and the background image to evaluate some measure of similarity between them. These measures include normalized cross correlation [11, 14], edge-width information and illumination ratio [15]. There are many other shadow detection methods as enumerated in [16].

1.1.3 Blob Extraction: This stage applies a connected component algorithm to the foreground mask to detect the foreground blobs while optionally discarding too small blobs created due to noise. Most existing methods use an efficient linear time algorithm that was developed in [17]. The popularity of this method is owing to the fact that it requires only a single pass over the image to identify and label all the connected components therein, as opposed to most other methods that require two passes [18, 19, 20].

1.1.4 Blob Tracking: This is often the most critical step in the AOD process and is concerned with finding a correspondence between the current foreground blobs and the existing tracked blobs from the previous frame (if any). The results of this step are sometimes used as feedback to improve the results of background modeling. Many methods exist for carrying out this task, including finite state machines [13], color histogram ratios [21], Markov chain Monte Carlo model [22, 23], Bayesian inference [24, 25], Hidden Markov models [26] and Kalman filters [27].

1.1.5 Abandonment Analysis: This step classifies a static blob detected by the tracking step as either abandoned or removed object or even a very still person. An alarm is raised if a detected abandoned/removed object remains in the scene for a certain amount of time, as specified by the user. The task of distinguishing between removed and abandoned objects is generally carried out by calculating the degree of agreement between the current frame and the background frame around the object's edges, under the assumption that the image without any object would show

better agreement with the immediate surroundings. There exist several ways to calculate this degree of agreement; two of the popular methods are based on edge energy [28, 11] and region growing [29]. There also exist methods [13] that use human tracking to look for the object's owner and evaluate the owner's activities around the dropping point to decide whether the object is abandoned or removed.

1.2 Analysis of previous research in this area: A great deal of research has been carried out in the area of AOD owing to its significance in anti-terrorism measures. Most methods developed recently can be classified into two major groups: those that employ background modeling and those that rely on tracking based detection. Some of the methods in the first group have been presented in [9, 10, 29-35]. Most of these use Gaussian Mixture Model (GMM) [3] for background subtraction. In this model, the intensity at each pixel is modeled as the weighted sum of multiple Gaussian probability distributions, with separate distributions representing the background and the foreground. The method used in [30] first detects blobs from the foreground using pixel variance thresholds and then calculates several features for these blobs to decrease false positives. The approach in [10] maintains two separate backgrounds- one each for long and short term durations- and modifies them using Bayesian learning. These are then compared with each frame to estimate dual foregrounds. The method detailed in [31] mainly focuses on tracking an object and its owner in an indoor environment with the aim of informing the owner if someone else takes that object. The method proposed in [29] applies GMM with three distributions for background modeling and uses these to categorize the foreground into moving, abandoned and removed objects. A similar background modeling method has been used in [34] along with crowd filtering to isolate the moving pedestrians in the foreground from the crowd by the use of vertical line scanning. There are also some approaches to background modeling that do not employ GMM such as the method in [9] that uses approximate median model for this purpose. Just like in [10], this one too maintains two separate backgrounds, one of which is updated more frequently than the other.

Some of the approaches based on the other class of methods, based on tracking, can be found in [21, 22, 36-39]. The tracking based approach used in [36] comprises three levels of processing- starting with background modeling in the lowest level using feedback from higher levels, followed by person and object tracking in the middle level and finally the person-object split in

the highest level to classify an object as abandoned. The system proposed in [37] considers the abandonment of an object to comprise of four sub-events, from the arrival of the owner with the object to his departure without it. Whenever the system detects any unattended object, it traces back in time to identify the person who brought it into the scene and thus identifies the owner. Tracking and detection of carried objects is performed using histograms in [21] where the missing colors in ratio histogram between the frames with and without the object are used to identify the abandoned object. The method used in [22] performs tracking through a trans-dimensional Markov Chain Monte Carlo model suitable for tracking generic blobs and thus incapable of distinguishing between humans and other objects as the subject of tracking. The output of this tracking system therefore needs to be subjected to further processing before luggage can be identified and labeled as abandoned.

1.3 Problem definition and scope: The problem that this work attempts to solve is concerned with the tracking and detection of suspicious objects in surveillance videos of large public areas. A suspicious object here is defined as one that is carried into the scene by a person and left behind while the person exits the scene. To be classified as abandoned, such an object should remain stationary in the scene for a certain period of time without any second party showing any apparent interest in it. In addition to detecting abandoned objects, this system also detects removed objects as any objects that were in the scene long enough to become part of the background and were subsequently removed.

The scope of this task is to identify any such suspicious objects in real time by looking for certain pre-defined patterns in the incoming video stream so as to raise an alarm without requiring any human intervention. It is assumed that the data about the scene is available from only one camera and from a fixed viewpoint.

The objectives of this system can be summarized as follows:

- It should be able to identify abandoned objects in real time and therefore must employ efficient and computationally inexpensive algorithms.
- It should be robust against illumination changes, cluttered backgrounds, occlusions, ghost effects and rapidly varying scenes.

- It should try to maximize the detection rate while at the same time minimizing false positives.

1.4 Formulation of the present problem: The overall problem of AOD can be broken down into a set of smaller, independent problems each of which is solved by a separate stage in the automatic AOD system. These stages are typically executed one after another with each stage using the output of the last stage as its input. There are five main stages in this process:

1.4.1 Background Modeling and Subtraction: This stage creates a dynamic model of the scene background and subtracts it from each incoming frame to detect the current foreground regions. The output of this stage is usually a mask depicting pixels in the current frame that do not match the current background model.

1.4.2 Foreground Analysis: The BGS step is often unable to adapt to sudden changes in the scene (of lighting, etc.) since the background model is typically updated slowly. It might also confuse parts of a foreground object as background if their appearance happens to be similar to the corresponding background, thus causing a single object to be split into multiple foreground blobs. In addition, certain foreground areas, while being detected correctly, are not of interest for further processing. The above factors necessitate an additional refinement stage to remove both false foreground regions, caused by factors like background state changes and lighting variations, as well as correct but uninteresting foreground areas like shadows.

1.4.3 Blob Extraction: This stage applies a connected component algorithm to the foreground mask to detect the foreground blobs while optionally discarding too small blobs created due to noise.

1.4.4 Blob Tracking: This is often the most critical step in the AOD process and is concerned with finding a correspondence between the current foreground blobs and the existing tracked blobs from the previous frame (if any). The results of this step are sometimes used as feedback to improve the results of background modeling.

1.4.5 Abandonment Analysis: This step classifies a static blob detected by the tracking step as either abandoned or removed object or even a very still person. An alarm is raised if a detected abandoned object remains in the scene for a certain amount of time, as defined by the user.

1.5 Organization of the thesis: Rest of this thesis is organized as follows: section 2 presents a brief description of the hardware and software used for developing and testing this application; section 3 describes the system methodology; section 4 describes the different modules in the system and their user interface; section 5 presents the details of testing methodology and the obtained results; finally, section 6 presents the conclusions and scope for future work.

2. Description of Hardware and Software Used

2.1 Hardware

Like most computer vision and image processing tasks, AOD is an extremely computationally intensive process and requires powerful hardware to run in real time. The present system has been tested on a fairly modern and moderately powerful laptop computer whose configuration is given below:

- *CPU*: Intel Core i5 3210M with 2 cores/4 threads running at 3.10 GHz
- *RAM*: 4GB DDR3 running at 1600MHz
- *GPU*: Nvidia Gefore GT635M with 96 CUDA cores clocked at 660MHz and having 1 GB of GDDR3 video memory clocked at 1800MHz
- *HDD*: 1 TB SATA-3 HDD with 5400 rpm

Testing has been done on both online and offline video streams. The online streams have been captured using both the integrated laptop camera and a dedicated Microsoft web-camera with respective resolutions of 1280x720 and 640x480. The offline test set includes videos from 3 publicly available benchmark datasets: PETS2006 [48], PETS2007 [49] and AVSS-iLiDS [50]. It also includes a custom dataset consisting of 18 videos shot using a Sony DSC-W35 digital camera. The video resolutions are 720x576 for public benchmark datasets and 640x480 for the custom prepared dataset.

2.2 Software

The application has been coded entirely in C++ and makes use of OpenCV library version 2.4.4. Coding has been done using Microsoft Visual Studio 2008 integrated development environment and tested on both Windows XP 32-bit and Windows 7 64-bit operating systems.

3. Theoretical Tools – Analysis and Development

This system considers the overall problem of AOD as a collection of smaller problems which can each be solved independently of the others. Thus it is built using a different subsystem for solving each of these sub-problems. Similar to the formulation stated in section 1.4, it can be divided into five main modules along with two additional modules for pre processing and for filtering the final results. The interactions between these modules are shown in Fig. 3 and a brief description of the algorithms that have been implemented for each of them follows.

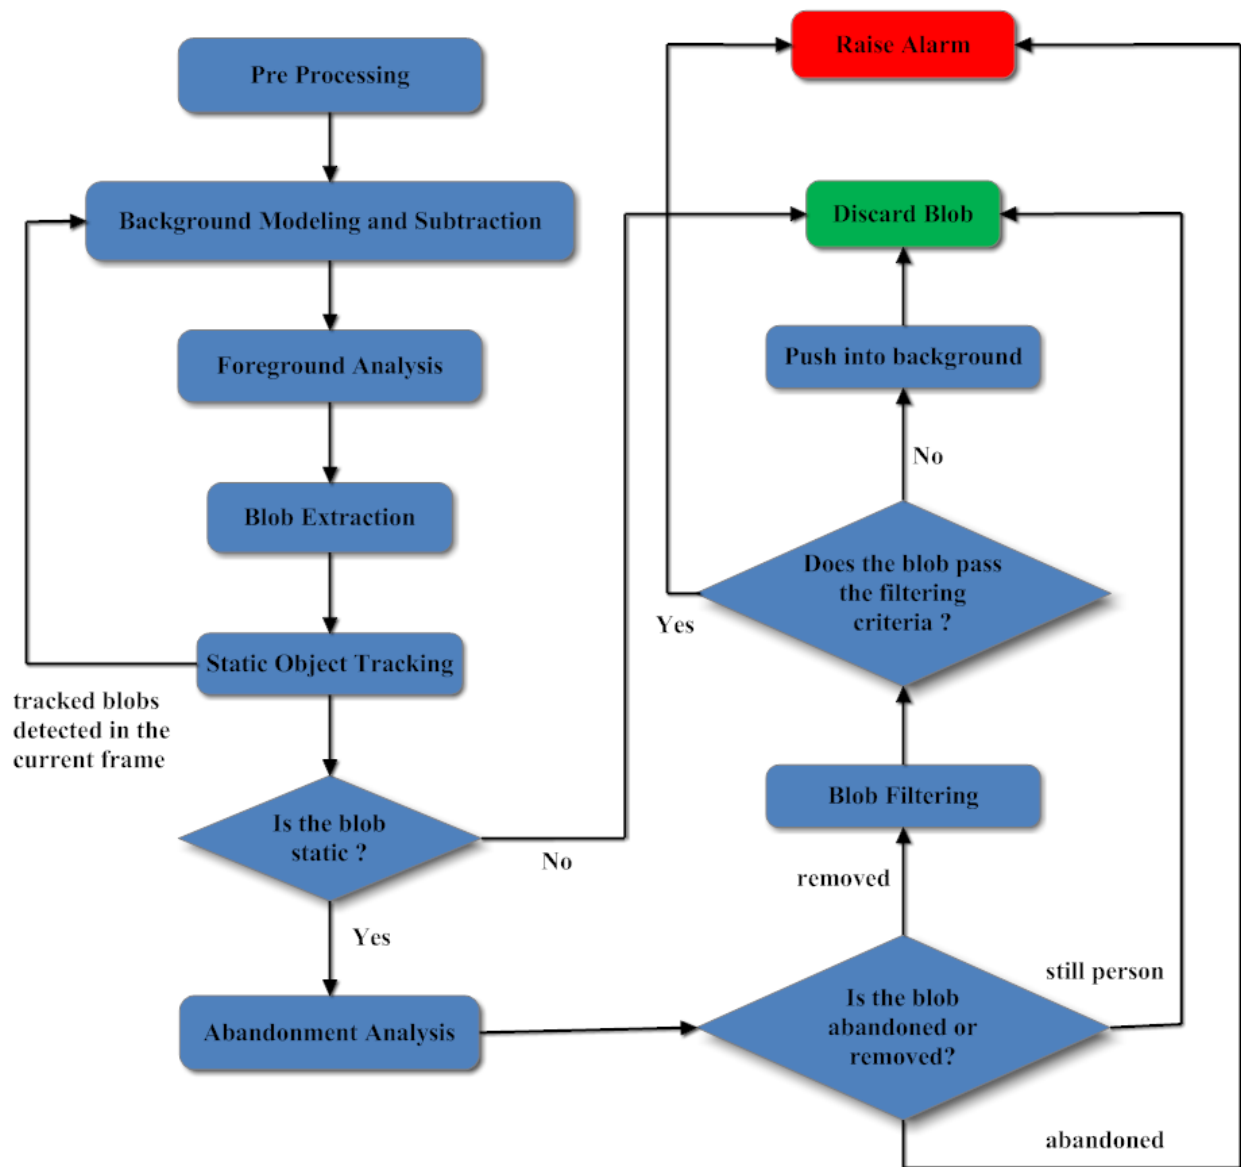


Fig.1 Flowchart of system methodology

3.1 Pre-processing: This module performs the following two functions:

3.1.1 Contrast enhancement: This step helps to improve the quality of low light videos like those taken at night (Fig. 1) by normalizing the difference between maximum and minimum intensities in the image which in turn helps to increase visibility in darker areas of the scene. Following three methods have been implemented for carrying it out:

3.1.1.1 Histogram equalization: This method involves the following steps:

- i. Split the input RGB image into 3 grayscale images, one corresponding to each channel.
- ii. Compute the histogram for each of these grayscale images and normalize it so that the sum of histogram bins becomes 255.
- iii. Compute the image corresponding to the transformed histograms using their integral.
- iv. Join the 3 transformed grayscale images to get the output RGB image.

Fig. 2 shows the input and processed images together with their respective histograms.

3.1.1.2 Linear contrast stretching: This method involves the following steps:

- i. Split the input RGB image into 3 grayscale images, one corresponding to each channel.
- ii. For each of these images, calculate the maximum and minimum cutoff intensities based on a user specified cutoff percentile and the image intensity distribution.
- iii. Set the intensity values of all pixels above the maximum and below the minimum cutoff as maximum and minimum intensity respectively.
- iv. Normalize (or stretch) the intensity distribution so that the minimum and maximum intensities become 0 and 255 respectively.
- v. Combine the transformed grayscale images to get the output RGB image.

Fig. 3 shows the result of applying this method.

3.1.1.3 Image filtering: This method involves the following steps:

- i. Split the input RGB image into 3 grayscale images, one corresponding to each channel
- ii. Apply the following formula to each pixel in each of these images:

$$I_{new}(i, j) = 5 * I_{old}(i, j) - [I_{old}(i - 1, j) + I_{old}(i + 1, j) + I_{old}(i, j - 1) + I_{old}(i, j + 1)] \quad (1)$$

- iii. Combine the transformed grayscale images to get the output RGB image.

Fig. 3 shows the result of applying this method.



Fig. 2 Result of applying histogram equalization on a dark scene from the custom dataset. The original image and processed images are at the top (left and right) while the corresponding histograms are at the bottom.



Fig. 3 Result of applying linear contrast stretching on a frame from AVSS PV-Night video. The original image and processed images are at the top (left and right) while the corresponding histograms are at the bottom.




Fig. 4 Result of applying image filtering on an indoor low-light scene from the custom dataset. The original image and processed images are at the top (left and right) while the corresponding histograms are at the bottom.

3.1.2 Noise reduction: This step reduces the white noise present in an input frame by smoothing the frame. It is particularly useful for low quality or grainy videos (Fig. 5) and is also needed to control the amount of noise that becomes visible in low light videos after applying contrast enhancement (Fig. 2,3). This is accomplished by subjecting the image one of several smoothing filters including linear, Gaussian and median filters. The result of applying Gaussian blurring filter is shown in Fig. 5.



Fig. 5 Result of applying linear convolution with a Gaussian kernel of size 5x5 on a frame from AVSS PV-Night video with the original image on the left and processed one on the right.

3.2 Background modeling and subtraction (BGS): The system currently includes three different BGS algorithms which are modified to perform object level background updating rather than the typical pixel level one. This is accomplished through a mask of currently tracked objects that is fed back from the blob tracking module (section 3.5) and is used to disable background updating at the corresponding pixels. This is needed to prevent static foreground objects from being learned into the background before they can be classified as abandoned or removed by the abandonment analysis module (section 3.6). A brief description of the three algorithms is presented below:

3.2.1 Gaussian mixture model (GMM): This method, first introduced in [3] and improved significantly in [4], models the distribution of pixel intensity values over time as a weighted sum of three Gaussian distributions. It assumes that the overall intensity at any pixel at each instant is produced by a combination of background and foreground processes, and each such process can be modeled by a single Gaussian probability distribution function. For each pixel in the current frame, the probability of observing the current intensity is given by:

$$P(X_t) = \sum_{i=1}^K \omega_{i,t} * \eta(X_t, \mu_{i,t}, \Sigma_{i,t}) \quad (2)$$

Here, K is the no. of distributions ($K=3$ here); $\omega_{i,t}$ is the weight associated with the i^{th} distribution at time t while $\mu_{i,t}$ is the mean and $\Sigma_{i,t}$ is the co-variance matrix of this distribution, η is the exponential Gaussian probability density function given by:

$$\eta(X_t, \mu_t, \Sigma_t) = \frac{1}{(2\pi)^{\frac{n}{2}} |\Sigma_t|^{\frac{1}{2}}} e^{-\frac{1}{2}(X_t - \mu_t)^T \Sigma_t^{-1} (X_t - \mu_t)} \quad (3)$$

Here, n is the dimensionality of each pixel's intensity value (e.g. $n=1$ for grayscale image and $n=3$ for RGB image). In order to avoid a costly matrix inversion and decrease computation cost, it is also assumed that the red, green and blue channels in the input images are not only independent but also have the same variance $\sigma_{k,t}^2$ so that the covariance matrix becomes:

$$\Sigma_{k,t} = \sigma_{k,t}^2 \mathbf{I} \quad (4)$$

The K Gaussian distributions are always ordered in the decreasing order of their contribution to the current background model. This contribution is measured by the ratio ω/σ under the assumption that higher is the weight and lower is the variance of a distribution, more is the

likelihood that it represents the background process. This assumption stems from the reasonable supposition that the background process not only has the maximum influence over a pixel's observed intensities but also shows very little variations over time (since the background is by definition static and unchanging). In the faster version of this method, presented in [4], the distributions are ordered simply by their weights, thus getting rid of the task of calculating ω/σ and also simplifying the sorting procedure without any significant impact on performance. Both of the above variants of GMM have been implemented in this system.

For each pixel in an incoming frame, its intensity value is compared with the means of the existing distributions starting from the first one and a match is said to be obtained if its Euclidean distance from the mean is less than m standard deviations ($m=3$ is used here), i.e. it satisfies the following condition:

$$|I_t - \mu_{k,t}| < m * \sigma_{k,t} \quad (5)$$

Here, I_t is the pixel intensity while $\mu_{k,t}$ and $\sigma_{k,t}$ are the mean and standard deviation of the k^{th} distribution at time t .

Since the background model is dynamic, it needs to be updated with each frame. While the weights are updated for all distributions, the mean and variance are updated only for the matched distributions. Following are the standard update equations used for this purpose:

$$\mu_t = (1 - \rho)\mu_{t-1} + \rho X_t \quad (6)$$

$$\sigma_t^2 = (1 - \rho)\sigma_{t-1}^2 + \rho(X_t - \mu_t)^T(X_t - \mu_t) \quad (7)$$

$$\omega_{k,t} = (1 - \alpha)\omega_{k,t-1} + \alpha(M_{k,t}) \quad (8)$$

Here X_t is the pixel intensity while $M_{k,t}=1$ for matched distributions and 0 for unmatched ones while ρ and α are learning rates. In the current work ρ and α are related as:

$$\rho_{k,t} = \frac{\alpha}{\omega_{k,t}} \quad (9)$$

Thus, while α is fixed for all distributions ($\alpha = 0.001$ used here), ρ is smaller for higher weighted distributions. If none of the existing distributions match the current intensity, the least

probable distribution (i.e. with the smallest value of ω/σ) is replaced by a new distribution with a high initial variance, low prior weight and the new intensity value as its mean.

3.2.2 Adaptive median: This BGS method, described in [1], works under the assumption that the background is more likely to appear at any given pixel over a period of time than foreground objects, i.e. the past history of pixel intensity values is likely to contain maximum occurrences of the background intensity at the pixel location. This leads to the reasonable supposition that the pixel stays in the background for more than half the values in its history. The median of previous n frames can therefore be used as the background model.

Though this BGS method is relatively easy to perform from a computational standpoint, it has fairly high memory requirements since the previous n frames must be stored in a buffer at any given time and n must be quite large to get a good estimate of the actual background. This is why a simpler recursive version of this algorithm is more practically feasible. In this approach, the running estimate of the median is incremented by one if the current intensity is larger than the existing estimate and decremented by one if it is smaller. The estimate is left unchanged if it equals the current pixel value. Following is the update equation for the background model in this approach:

$$B_t = \begin{cases} (B_{t-1} - 1) & \text{if } B_{t-1} > I_t \\ B_{t-1} & \text{if } B_{t-1} = I_t \\ (B_{t-1} + 1) & \text{if } B_{t-1} < I_t \end{cases} \quad (10)$$

Here, B_t and I_t respectively refer to the intensity values in background model and the current frame at time t . If run over a sufficiently long period of time, this running estimate would eventually converge to a value that is larger than half the pixel values encountered so far and smaller than the other half and is therefore a good approximation to the median.

3.2.3 Running Gaussian average: The basic idea here is similar to that in the last method except that here the average, rather than the median, of the last n frames is used as the background model. Using a non recursive approach here too is computationally inexpensive but memory consuming thus leading to the use of running averages.

Since recent frames are more likely to contribute to the current background than older ones, a weighted average is used with higher weights attached to more recent frames. When these

weights vary according to the Gaussian distribution, the running Gaussian average is obtained. This process can alternatively be interpreted as the fitting of a single Gaussian distribution over the pixel intensity histogram.

The algorithm implemented in this system is detailed in [2]. The background model here is updated according to the following equations:

$$\mu_{t+1} = \alpha I_t + (1 - \alpha)\mu_t \quad (11)$$

$$\sigma^2_{t+1} = \alpha(I_t - \mu_t)^2 + (1 - \alpha)\sigma^2_t \quad (12)$$

Here, μ_t and σ^2_t respectively refer to the mean and variance of the Single Gaussian distribution while α is the learning rate ($\alpha = 0.001$ is used in this work).

3.3 Foreground analysis: Most existing BGS methods are far from perfect and tend to produce noisy outputs that need to be processed further before they can be used to extract useful objects. The noisy portions of the BGS output may contain false foregrounds produced for example during sudden lighting changes (a light is switched on or off) as well as actual foregrounds that are of no interest to further processing but can complicate it significantly. Foreground regions detected due to moving shadows are important examples of this latter category. Thus, this system has a separate foreground analysis stage that takes as its input the noisy foreground mask produced by the BGS module and removes all the false and uninteresting foreground pixels from it. This process is further divided into the following three distinct tasks:

3.3.1 Detecting sudden lighting changes: When a new light source is suddenly introduced into a scene, it creates a patch of light that is detected as a foreground region since its color intensity values are significantly different from those of the existing background model that has not yet adapted for the new light source. However, though the actual RGB values may have changed, the underlying texture of the region remains unchanged and this fact can be utilized to detect such false foreground regions. A texture difference measure, proposed in [11], is used for this purpose. Let g_f and g_b denote the gradient vectors for the current and background frames respectively where $g_f = (g_f^x, g_f^y)$ and $g_b = (g_b^x, g_b^y)$. The partial derivatives g^x and g^y are calculated using Sobel operator with a neighborhood window size of 5x5. The texture difference at a particular pixel can then be defined as:

$$S = \frac{\sum_{u \in W} (2 \|g_f\| \cdot \|g_b\| \cos \theta)}{\sum_{u \in W} (\|g_f\|^2 + \|g_b\|^2)} \quad (13)$$

Here W represents a $M \times N$ window ($M = N = 5$ is used here) centered at that pixel while θ is the angle between the vectors g_f and g_b . In fact, the numerator in the above ratio can be conveniently calculated as $2(g_f^x * g_b^x + g_f^y * g_b^y)$ thus dispensing the need to calculate θ . For false foreground regions the ratio $S \approx 1$ while it is significantly less for actual foreground objects. Thus, a pixel is classified as false foreground if $S \geq T_{frg}$ where $T_{frg} = 0.7$ has been used in this system.

Though the above method is quite good at detecting localized lighting changes, it does not work for scene wide changes (e.g. lights being switched on/off in a room as in Fig. 4). However, in such cases, an unusually large part of the scene is detected as foreground and this fact can be used to detect such changes. This system continuously monitors the fraction of frame pixels classified as foreground and if this fraction exceeds a threshold, say 0.60, for several consecutive seconds, the system resets the background model with the current frame and thus compensates for the lighting change.

3.3.2 Detecting shadows: Two different shadow detection methods, both utilizing the normalized cross-correlation (NCC) of intensity values between the grayscale current frame and background image, have been implemented in this system. The two methods, though similar, differ in the actual formulation of NCC. They are described below:

3.3.2.1 Complex NCC: This is a single step process described in [11]. The NCC is calculated for each pixel according to the following expression:

$$NCC = \frac{\sum_{u \in W} I_f \cdot I_b - \frac{1}{MN} \sum_{u \in W} I_f \sum_{u \in W} I_b}{\sqrt{(\sum_{u \in W} I_f^2 - \frac{1}{MN} (\sum_{u \in W} I_f)^2) (\sum_{u \in W} I_b^2 - \frac{1}{MN} (\sum_{u \in W} I_b)^2)}} \quad (14)$$

Here, W denotes the $M \times N$ neighborhood centered at that pixel while I_f and I_b respectively denote the current frame and background intensity values at a particular pixel. A pixel is classified as shadow if $NCC \geq T_{shadow}$ and $I_f \geq T_{intensity}$. The second condition has been added

to avoid misclassification of very dark areas as shadows. $T_{shadow} = 0.60$ and $T_{intensity} = 5$ have been used in this work.

3.3.2.2 Simple NCC: This is a two step process described in [14]. The first step, that utilizes NCC, is to identify candidates for shadow detection while the second step, that uses neighborhood statistics, is the shadow refinement stage to prevent misclassifications. Unlike the first method, this one uses a much simpler expression for NCC:

$$NCC = \frac{\sum_{u \in W} (I_f \cdot I_b)}{\sqrt{(\sum_{u \in W} I_f^2)(\sum_{u \in W} I_b^2)}} \quad (15)$$

The different symbols here have the same meanings as in the previous method. NCC is close to unity for shadowed regions. Since a shadowed region is expected to be darker than a non-shadowed one, the intensity energy in a pixel's neighborhood must be less in the current frame.

Intensity energy here is defined as $E_f = \sqrt{\sum_{u \in W} I_f^2}$ for the current frame and $E_b = \sqrt{\sum_{u \in W} I_b^2}$ for the background. Thus the two conditions to be satisfied for pre-classification as shadow are $NCC \geq T_{NCC}$ and $E_f < E_b$.

Once a pixel is detected as a shadow candidate, it is subjected to a refinement stage that checks whether the ratio I_f/I_b is relatively constant in a neighborhood around the pixel. This is checked

by calculating the standard deviation of I_f/I_b within the neighborhood. To make allowance for

very dark pixels that may have an intensity of zero, the current system evaluates $I_f + 1/I_b + 1$ instead. Just like in the last method, this one too includes the provision of

avoiding the misclassification of very dark pixels as shadows through an additional constraint on the permissible range of values of this ratio. Thus the two conditions to be satisfied to confirm a pixel as shadow are $std(I_f + 1/I_b + 1) < T_{std}$ and $T_{min} \leq I_f + 1/I_b + 1 < 1$. The second

condition here must be satisfied by all the pixels in the neighborhood. Values of $T_{NCC} = 0.95$,

$T_{std} = 0.05$ and $T_{min} = 0.5$ have been used in this system. Neighborhood size of 3x3 has been used for both shadow detection methods.

Fig. 6 shows an example of shadow removal in the foreground mask.

3.3.3 Morphological frame processing: Both lighting change detection and shadow detection processes are prone to false classifications and often leave behind ‘holes’ inside valid foreground objects along with some left behind shadow pixels that may be detected as small blobs. These are removed by applying the morphological operations of closing followed by opening. The former process applies dilation followed by erosion by while the latter applies erosion followed by dilation. An example of the cleaning-up effect of morphological processing is shown in Fig. 7.




Fig. 6 Result of applying foreground processing (bottom right) on the foreground mask (top right) obtained using the current frame (top left) and the current background (bottom left).




Fig. 7 Foreground mask before (left) and after (right) applying morphological processing.

3.4 Blob extraction: The refined foreground mask produced by the last step is subjected to a connected component detection algorithm to extract meaningful foreground objects while discarding any blobs that are smaller than a specified threshold. A simple, efficient but fairly accurate algorithm described in [17] is used for this purpose. This algorithm uses a single pass over the image to detect external as well as internal contours and also labels the internal points of each connected component.

3.5 Blob tracking: This stage tracks only the static objects in the scene using a heavily modified version of the algorithm used in [9]. Following are the main steps in this process:

3.5.1 Establish blob correspondence: The first step is to compare each blob in the incoming frame with the existing blobs in the tracking system to find a match based on position and size. For an existing blob to match a new blob, the difference between their positions, as measured by some distance measure, must be less than a user specified fraction of the length of its bounding box diagonal. Two distance measures have been used here: one is the Euclidean distance between the centroids of the two blobs while other is the Euclidean distance between one blob's centroid and the nearest point in the other one's bounding box. If the positions of two blobs are found to match, they are next compared by their sizes, measured either by the contour areas or the bounding box areas. These must again differ by less than a user specified threshold for a final match to occur.

An additional constraint is imposed on the matching process that each existing blob may match at most one new blob and vice versa. All new blobs that do not match any existing blob are added to the tracking system with their hit counts initialized to one.

3.5.2 Update state variables for existing blobs: Three state variables are maintained for tracked blobs: hit count, miss count and occluded count. When a match is found for an existing blob, its hit count is incremented by one while its miss and occluded counts are set to zero.

An existing blob that does not match any of the new blobs is checked for occlusion by assuming that if more than a certain fraction (say 0.8) of the blob's pixels has been detected as foreground in the current frame, then it is occluded by one or more foreground objects. In this case, both its hit count and occluded count are incremented by one. If the object is not detected as occluded, it is considered as missing in the current frame and its miss count is incremented by one.

3.5.3 Remove non-static objects: An object is discarded from the tracking system if one of the following conditions is satisfied:

- Not detected/matched for several consecutive frames: This occurs when its miss count exceeds either a user specified threshold or the current hit count.
- Occluded for too long: This occurs when the occluded count exceeds a threshold.
- Appearance changes significantly.

A blob's appearance information is maintained by storing the incoming frame (or its gradient) when the blob is first added to the tracking system. If the object is indeed static, its appearance would not change significantly in subsequent frames as long as it remains visible (i.e. not occluded). This change in appearance is measured by the mean difference in pixel intensity between the stored image and the current frame averaged over all the pixels in the blob's bounding box. If this difference exceeds a threshold, the object is immediately discarded. An exponential moving average of these mean differences is also maintained for use by the abandonment analysis module for distinguishing between actual static objects and very still persons.

3.5.4 Change object label: If the hit count of an object exceeds a user defined threshold, it is labeled as static and its maximum permissible occluded and miss counts are multiplied by

respective constant factors to make greater allowance for accidental misses and prolonged occlusions. If it exceeds a second higher threshold, this object is passed to the abandonment analysis module to be classified as abandoned or removed.

3.5.5 Provide feedback to the BGS module: To prevent static objects from being learnt into the background, all the tracked objects that are detected in the current frame fed back to the BGS module so that the background model is not updated for pixels contained in their bounding boxes. An object as a whole is pushed into the background if one of the following conditions hold true:

- i. It is detected as a state change in the abandonment analysis module.
- ii. It is detected as abandoned or removed and the alarm timeout has been reached.

This feedback enables the BGS algorithm to perform object-level (as opposed to pixel-level) background updating and thus helps to avoid foreground fragmentation in addition to maintain the objects of interest in the foreground.

3.6 Abandonment analysis: This is the final stage whose purpose is to prevent false detections of abandoned objects due to the ‘ghost effect’ created when an object in the background is removed, leaving behind a false foreground object. It also checks the blob’s internal variation during the period it has been tracked to ensure that it is not a very still person. This module therefore can be divided into two sub-modules:

3.6.1 Removed object detection: When a background object is removed from the scene, it is reasonable to assume that the area thus vacated will now exhibit a higher degree of agreement with its immediate surroundings than before since that object was presumably hiding a part of the relatively uniform backdrop that was visible in its neighborhood. Conversely, when a new object is added to the scene, the region where it lies is likely to show a decreased level of conformity with its surroundings. This assumption can be used to classify an object as abandoned or removed by evaluating, in both the background image and the current frame, a measure of agreement of the object’s edge (and nearby interior) pixels with their immediate neighborhood pixels (outside the object). Provided that the two images exhibit a significant difference in conformity levels, the object is likely to be an abandoned object if the background image shows greater agreement and a removed object otherwise. If the two images exhibit fairly

similar degrees of conformity, the object is likely to represent a state change (like a door closing). Following two methods have been implemented to measure this degree of agreement in this system:

3.6.1.1 Region Growing: This method was introduced in [29] and consists of a two step process. In the first step, the blob's boundary is eroded to obtain a set of pixels that lie near the boundary but are completely inside the object. These pixels are then used as seed points in a standard region growing algorithm to add all the surrounding pixels that are similar to these. This process is repeated for both the background and the current frame. The image not containing the object is likely to exhibit a significantly greater number of added pixels than the other one where the region growing process would stop around the object's actual boundaries. Thus, the object is considered removed if the region growth is significantly more in the current frame and abandoned otherwise.

3.6.1.2 Edge Detection: This method, used in [11, 28] works on the idea that placing an object in front of the background will introduce more edges to the scene around the object's boundaries, provided that the background is not extremely cluttered. Two methods have been implemented for edge detection: one is a simple gradient based method that uses Sobel operator to obtain separate x and y gradient images while the other one uses Canny edge detection [46]. The results obtained using both these methods are shown in Fig. 8. The points located along the object's external contours are checked for high intensity values (that indicate edges) in the gradient images obtain the edge energy associated with that object. The object is classified as removed if this energy comes out to be greater for the background image and abandoned otherwise.




Fig. 8 Results of applying gradient-based (center) and Canny (right) edge detection on the same frame (left).

3.6.2 Still person detection: As mentioned in section 5, an exponential running average of mean pixel differences between the current frame and the stored image is maintained for each object in the tracking system. The basic idea here, introduced in [45] is that while the appearance of a true static object remains completely unchanged (provided that it is not occluded) frame after frame while a false static object detected due to a very still person will show some variations in the pixel values within the constant bounding box (unless the person is standing perfectly and unnaturally still). These internal blob variations are captured by the running average of mean differences and if this value exceeds a certain threshold, the object is classified as a still person. In order to make this process immune to lighting changes, gradient images, instead of the actual images, can be used to calculate these mean differences.

3.7 Blob filtering: The abandonment module classifies a static blob into one of four categories: abandoned, removed, state change or still person. A blob detected as a state change or a still person is removed from the tracking system. If a state change, it is also pushed into the background. An object detected as abandoned causes an alarm to be raised immediately while one detected as removed is first passed through a filtering process where it is compared to each of the user-specified regions in the scene (provided that any have been added to the filtering module). An alarm is raised only if it matches one of them. This matching can be done on the basis of the blob's position, size, appearance or any combination of these, as specified by the user.

4. Development of Software

The present system has been built using a modular approach with several distinct sub-systems, each of which solves a particular aspect of the AOD problem. Since this system is meant to be more of a flexible framework than a stand-alone system, multiple methods have been implemented for each of these sub-systems with the possibility of adding more in the future without too much re-programming effort. All of these methods have been integrated into the system in such a way that the user has the option to not only tweak the parameters of any one method but to switch between different methods in real time while observing its impact in terms of both accuracy and speed. It is also possible to switch between the online and offline input sources without having to restart the system run.

The application's user interface has been designed to allow the user to interact with the system in the following ways:

- Adjust each of the system parameters and switch between the different methods implemented in each module in real time. This can help to quickly find the optimal values of various parameters and determine the combination of methods that provide the best performance in any given scenario.
- Switch between different input sources (camera or video file) as well as different datasets (or different videos in the same dataset) in real time without having to restart the system run. This can be useful for gauging the performance of a particular combination of methods and parameters under different scenarios.
- Selectively modify the BGS process by selecting some or all of the blobs in the current frame to be pushed into the background. This can help the system to take advantage of the vastly superior human visual processing ability to improve its performance in very difficult scenarios while still being completely autonomous most of the time.
- This feature is particularly useful at the beginning of the video stream processing (or after a scene wide appearance change) when the background model differs significantly from the true background, thus leading to several false foreground blobs. These can not only slow down the processing speed but can also occlude actual foreground objects of interest. The user, in such cases, can conveniently push all unwanted blobs into the background at the press of a key.
- Specify particular locations in the scene which are to be monitored for removed objects (refer section 3.7). For example, the user may want an alarm to be raised only for some valuable items in the scene (e.g. a CPU or monitor in a computer lab) but not for non-important items (e.g. chairs).
- Alter the input video stream either by pausing or by fast-forwarding it to a particular frame. The former can be used to test the results of applying different processing methods (e.g. detecting edges using either Sobel gradient operator or Canny edge detection) on the

same frame while the latter can be useful for quickly getting to the point of interest (e.g. the actual drop or removal event).

- Adjust the processing resolution in real time and observe its impact on both performance and accuracy.

This application has been written entirely in C++ using OpenCV library that allows track bars and mouse/keyboard events as the only options for user interactions. Therefore parameter adjusting and switching between different methods, input sources or datasets are all accomplished using track bars. Each module in the system displays one or more windows where that module's output is displayed along with all the relevant track bars. A couple of sample windows are shown in Fig. 8. Keyboard events are used for modifying the BGS process while mouse events are used for selecting objects to remove (push into the background) or add (for filtering).

The system can take its input from a video file as well as from a camera attached to the system. The video files must have their filenames formatted in the certain manner that specifies their set, camera angle and size. The shows the output images of all the stages in different windows and also saves the important ones into separate video files so that they can be run later at the original frame rates.

An important aspect of this system is the large number of user-adjustable parameters that it offers. Each of the five main modules has its own parameters and in all there are currently more than 50 parameters, all of them adjustable in real time. The system reads the initial values of these parameters from an ASCII text file that can be formatted in a user friendly way with comments describing the purpose and possible values of each of these parameters. During the program execution, each parameter is assigned its own sliding track bar in the appropriate window and can be conveniently adjusted by moving the slider. Each of the subsystems creates one or more windows to display its output and place the track bars. Since OpenCV places a very strict limit on the length of track bar labels, the user is also provided a detailed description of each track bar along with its possible values whenever a mouse click is received anywhere inside the concerned window (Fig. 10,11,15).

Following are the main modules along with screenshots and descriptions of the interface they provide to the user:

4.1 User interface module: This module is responsible for allowing the user to change the input source in real time. It also allows the user to adjust the processing resolution in case it is desired to be lower than the original input resolution so as to increase the processing speed. It creates two windows: one shows the current input frame in its original resolution along with the input source parameters while the other one shows a combined image with current frame, processed and un-processed foreground masks, the background image and the results produced by the abandonment module (either edge detection or region growing) for both current and background images. Both of these windows also show the average and current frame rates obtained with the present configuration. The images displayed in the second window are all at the processing resolution. Examples of the two windows are shown in Fig. 9.

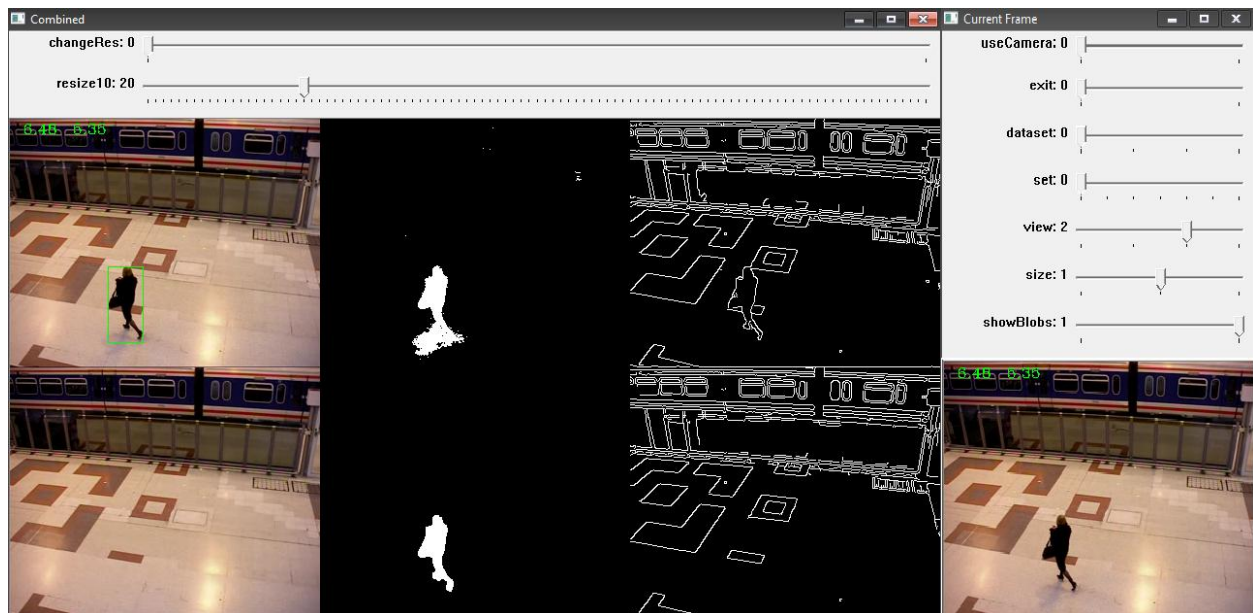


Fig. 9 Windows created by the User interface module showing the processing (left) and input source (right) parameters

4.2 Pre-processing module: As stated in section 3.1, this module carries out two functions: contrast enhancement and noise reduction. It creates one window that allows the user to enable or disable these functions, choose one of the several methods for carrying each out and adjust the parameters relevant to the chosen method. It also provides an option to display the histograms of images computed before and after they are pre-processed. The image displayed in this window is

the result of applying these functions on the current frame. Fig. 10 shows window created by this module along with the help associated with the same.

4.3 BGS module: This module utilizes an online OpenCV based BGS library [50] that contains generic implementations of 7 BGS methods out of which 4 have been used in this work. There are two GMM methods, one based on [3] and the other on [4]. Both of these have been heavily modified (practically rewritten) to accommodate the specific requirements of this system. The other two methods have been used relatively unchanged.

This module creates a single window that displays the current background image along with the parameters of the currently active BGS method. It also provides options to change the active BGS method and to adjust the parameters relevant to the chosen method. An example window is depicted in Fig. 11 along with its associated help.

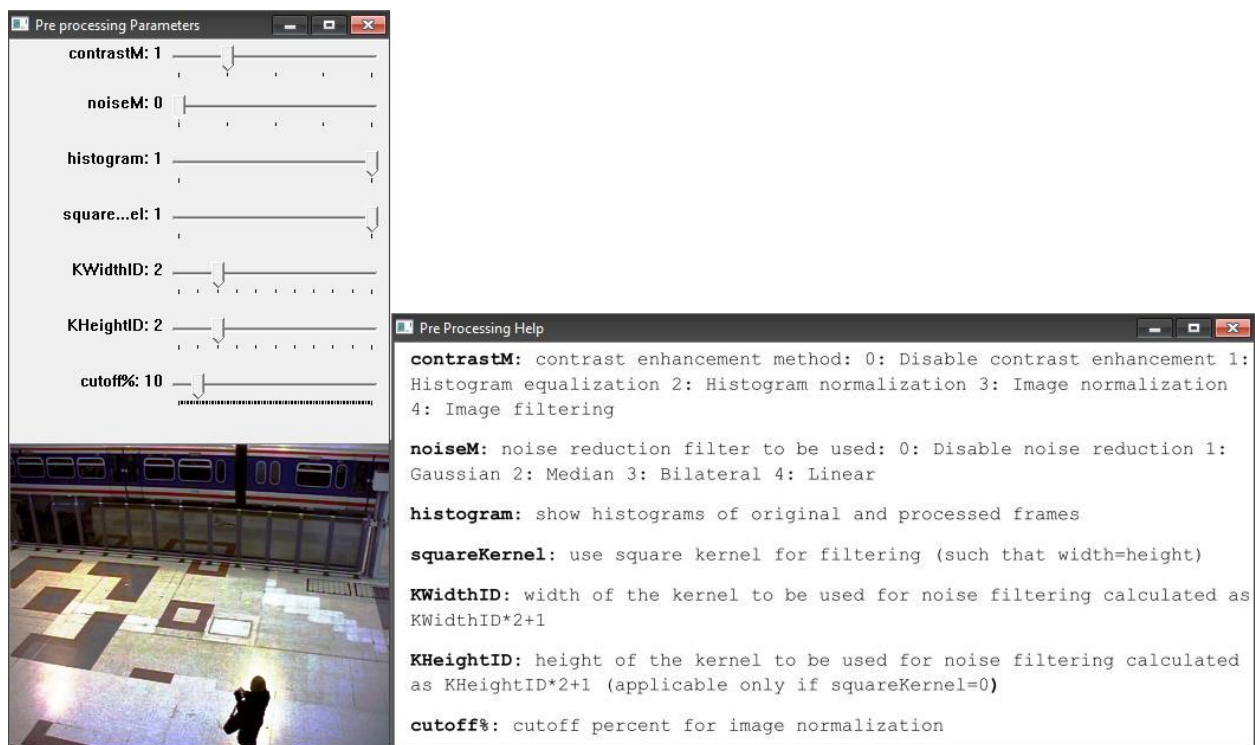


Fig. 10 Left: Window created by the pre-processing module showing the relevant options and the result of applying contrast enhancement. Right: Help associated with this window.

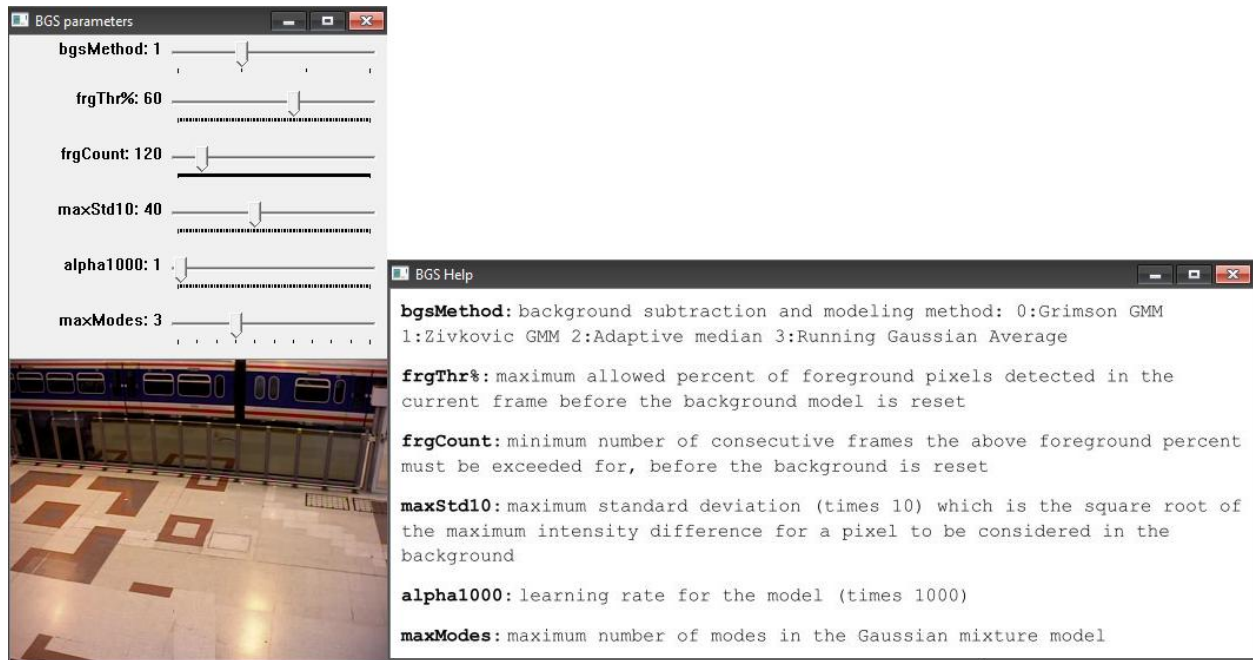


Fig. 11 Left: Window created by the BGS module showing the result of Zivkovic GMM BGS method and its relevant parameters; Right: The associated help.

4.4 Foreground analysis module: This one creates two windows: one for adjusting the various parameters of the false foreground and shadow detection process and other one for turning on or off various stages of the foreground analysis process. The first one displays the result of lighting change and shadow removal process while the second one shows the result of applying the morphological operations. Screenshots of these windows are shown in Fig. 12.

4.5 Blob extraction module: This one uses a modified version of an online blob extraction library called ‘cvblobslib’ [51]. It creates a single window that shows the blobs extracted from the current frame (Fig.13) along with the minimum blob size parameter that can be adjusted.

4.6 Blob tracking module: This module creates two windows: one to display the parameters that control the process of finding correspondence between new and existing blobs and the other one to show the parameters governing the tracking system. The second window also displays the current frame with the bounding boxes of all currently tracked blobs while the first one shows the current tracked object mask that is used as feedback to modify the BGS process (Fig. 14).

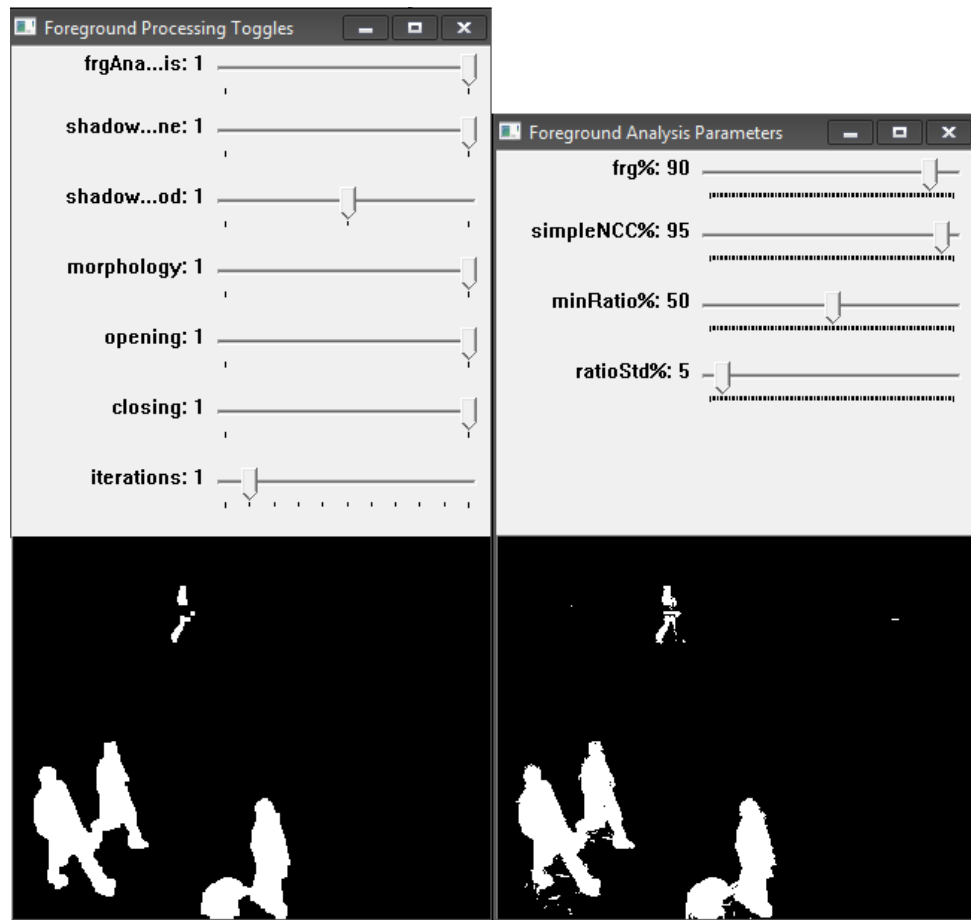


Fig. 12 Windows created by the foreground analysis module with the parameter window on the right and toggle window on the left.



Fig. 13 Windows created by the blob extraction module with the blobs extracted from the current frame shown in green.

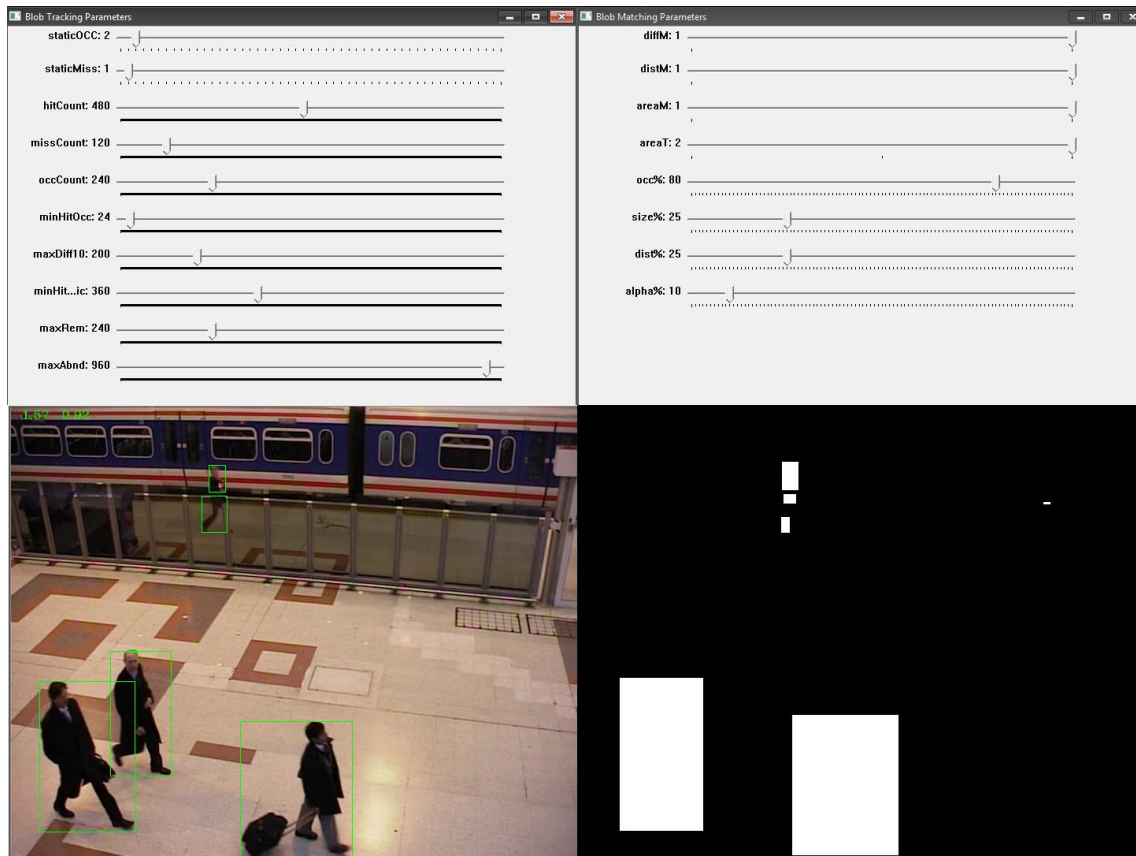


Fig. 14 Windows created by the blob tracking module with the tracking parameter window (showing currently tracked blobs) on the left and matching parameter window (showing the tracked object mask) on the right.

4.7 Abandonment analysis module: This module consists of two sub-modules: one for region growing and other one for edge energy based abandonment analysis. It creates a single window that shows the parameters of either region growing or edge detection based method, depending on which is currently active, along with the still person detection threshold. It also displays the result produced by the selected method (grown regions for the former and detected edges for the latter). Fig. 15 depicts a screenshot of this window next to its associated help window.

4.8 Blob filtering module: This one creates a single window that shows the objects currently added to the system by the user for filtering. It also provides options to toggle size, location and appearance based filtering together with all relevant parameters. When the user indicates that one or more objects are to be added it creates another window displaying the current frame where the user can use mouse clicks to specify bounding boxes of the objects to be added. Screenshots of both these windows are presented in Fig. 16.

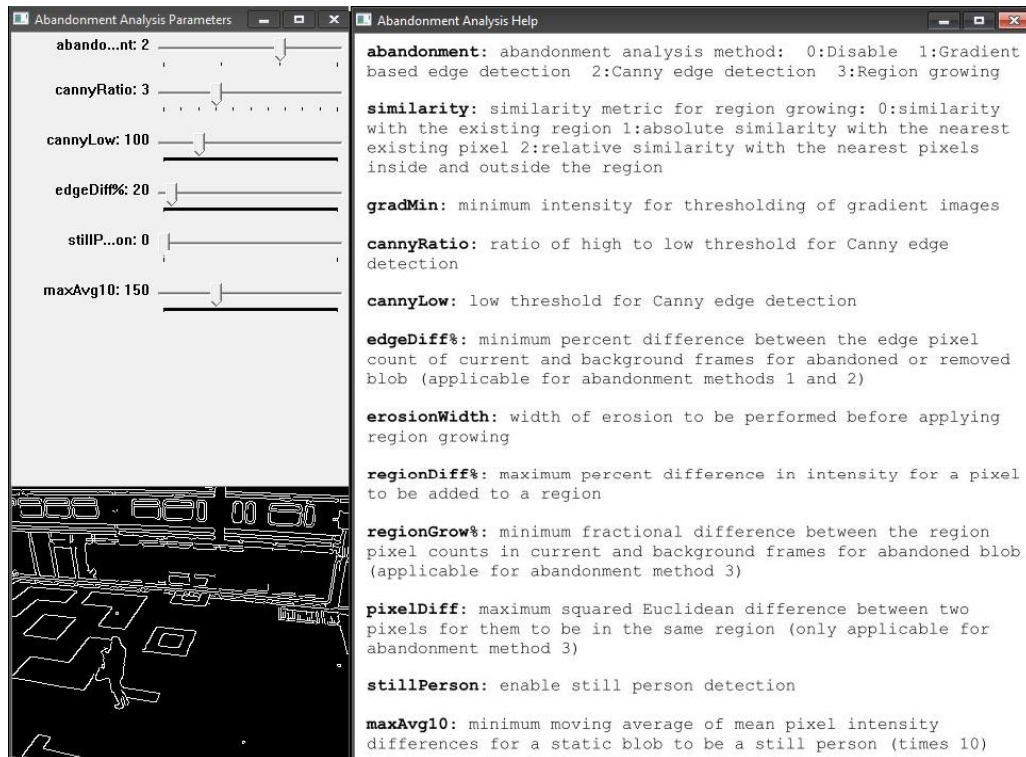


Fig. 15 Window created by the abandonment analysis module (left) with its associated help (right).

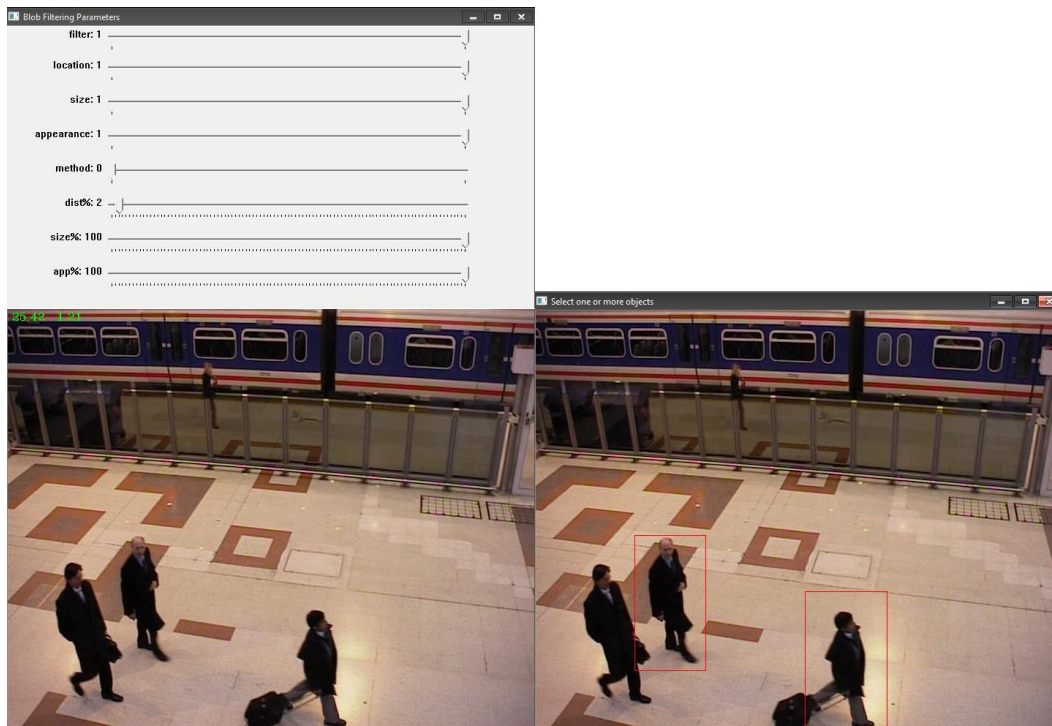


Fig. 16 Window created by the abandonment analysis module with the parameter window on the left (showing that no objects have been added yet) and the object selection window on the right where two bounding boxes (in red) have been drawn by the user to add them to the filtering system.

5. Testing and Analysis

5.1 Setup: Since the system features multiple methods for several of its modules, it is not possible to provide the results of every possible combination of these methods. Therefore, some preliminary testing was done and the following combination of methods was found to give good results in most cases:

- Pre-processing: Both contrast enhancement and noise reduction were turned off except in a few night/low light scenes.
- BGS: modified GMM introduced in [4] (section 3.2.1).
- Shadow detection: Simple NCC based method (section 3.3.2)
- Morphological processing: Both closing and opening operations were enabled (section 3.3.3).
- Abandonment analysis: Canny edge detection (sub-section 3.6.1.2).

The results presented here have been obtained using this configuration for all the tests with only a few minor modifications for specific cases.

5.2 Datasets: This system has been tested on 4 different databases, 3 of which are publicly available benchmark databases while one is a custom prepared dataset with videos of various indoor and outdoor scenarios featuring both day and night time scenes. Following is a brief description of these datasets:

5.2.1 PETS 2006: This dataset, available at [47], contains videos of 7 different events each captured from 4 different viewpoints. The videos depict various left-luggage scenarios of increasing complexity captured in a train station. A wide range of items including briefcase, suitcase, rucksack and even a ski gear carrier constitute the left-luggage in these videos. Fig. 17 depicts sample scenarios from this dataset from all 4 views.

5.2.2 PETS 2007: This dataset, downloadable from [48], contains videos of four scenarios: loitering, attended luggage removal, luggage exchange, and unattended (or abandoned) luggage. There are two events for each of these scenarios with each event having been captured from 4 angles. Only the last two sets of videos (sets 7 and 8) depict abandoned object events and hence only these have been used for testing. A few sample scenes from this dataset are presented in Fig. 18.




Fig. 17 Sample images from PETS2006 dataset




Fig. 18 Sample images from PETS2007 dataset

5.2.3 AVSS i-LIDS: This dataset is available at [49] and consists of CCTV footage of two scenarios: abandoned baggage (AB) and illegally parked vehicles (PV). There are 3 different videos of increasing difficulty levels for both of these scenarios with the latter also having a night time video. Testing was done on both scenarios since a parked vehicle can also be considered as an abandoned object. Fig. 19 shows both scenes from both AVSS-AB (top row and middle left) and AVSS-PV (bottom row and middle right) datasets.



Fig.19 Sample images from AVSS i-LIDS dataset.

5.2.4 Custom dataset: This dataset contains 18 videos out of which 12 are outdoor scenes while 6 are indoor ones; 13 of these were shot during in well lit conditions while 5 were shot at night or in poorly lit interiors. Most videos feature both abandoned and removed object events. Additional information about these videos, including the time and place of shooting and their durations is presented in Table 1. Fig. 20 depicts outdoor day (top left), outdoor night (top right), indoor dark (bottom left) and indoor well lit (bottom right) scenes from the custom dataset.

Videos in the first three datasets have original resolution of 720x576 but have been resized to 360x288 for testing. The custom dataset videos were shot at 640x480 and then resized to 320x240 for testing.



Fig.20 Sample images from the custom prepared dataset.

Set	Time/Lighting	Location	Duration (sec)
1	daytime	outdoor	131
2	daytime	outdoor	91
3	daytime	outdoor	84
4	daytime	outdoor	268
5	daytime	outdoor	375
6	daytime	outdoor	263
7	daytime	outdoor	260
8	daytime	outdoor	201
9	daytime	outdoor	171
10	night	outdoor	263
11	night	outdoor	310
12	night	outdoor	331
13	well lit	indoor	183
14	well lit	indoor	296
15	dark	indoor	346
16	well lit	indoor	535
17	well lit	indoor	188
18	dark	indoor	91

Table 1 Description of the custom dataset.

5.3 Result summary: Experiments were carried out on an Intel Core i5 processor clocked at 2.5 GHz and having 4 GB of RAM. Processing was carried out at both the original resolution (360x288 or 320x240) as well as at half the resolution in each dimension (180x144 or 160x120). Average processing speed of around 9 fps was obtained for the former and 20 fps for the latter. In most cases, half-resolution processing was sufficient to detect the suspicious events without giving any false positives; full resolution processing was required only in some of the more complex situations. The publicly available datasets have only been tested for abandonment

events while the custom dataset has been tested for both abandonment and removal events. The results have been summarized below in Tables 2 and 3.

Set		Actual events	True detections	False detections	
				Still Persons	Other objects
PETS2006					
1		4	4	0	1
2		4	4	0	0
3		0	0	0	0
4		0	0	1	0
5		4	4	1	1
6		4	4	0	1
7		4	3	0	1
Total		20	19	2	4
PETS2007					
7		4	3	0	0
8		4	3	0	0
Total		8	6	0	0
AVSS i-LIDS					
AB	Easy	1	1	0	0
	Medium	1	1	0	0
	Hard	1	1	1	1
	Total	3	3	1	0
PV	Easy	1	1	0	0
	Medium	1	0	0	1
	Hard	1	0	0	1
	Night	1	1	0	0
	Total	4	2	0	2
Total		7	5	1	3
Custom dataset					
1		1	1	0	0
2		1	1	0	0
3		0	0	0	0
4		2	2	0	0
5		2	2	0	0
6		2	1	0	0
7		2	0	0	0
8		2	2	0	0
9		1	0	0	1
10		2	2	0	0
11		2	2	0	0
12		2	2	0	0
13		2	2	0	0
14		2	2	1	0
15		2	0	0	0
16		8	8	0	1
17		1	1	0	0
18		1	0	0	0
Total		35	28	1	2
Overall		70	58	4	9

Table 2 Result summary for all the datasets

Category	Actual events	True detections	False detections	
			Still Persons	Other objects
Outdoor	19	15	0	1
Indoor	16	13	1	1
Daytime / Well lit	26	22	1	2
Night/Dark	9	6	0	0

Table 3 Category wise result summary for custom dataset.

5.4 Result analysis: A modified version of the quantitative metric introduced in [9] has been used here to measure the overall performance of this system. This metric evaluates the number of successfully detected events, penalized by the false detections, as a fraction of the total number of events. It is given by:

$$S = \frac{\text{true detections} - (0.5*n_p + 0.25*n_b)}{\text{total no.of events}} \quad (16)$$

Here, n_p and n_b respectively represent the no. of still persons and other objects falsely detected as abandoned or removed objects. This expression is based on the assumption that the negative impact on performance of detecting a still person as abandoned/removed object is half the positive impact due to a true detection while a false detection involving a non-person object has a quarter of the negative impact. The rationale behind this assumption is that, since the former scenario (a person standing very still) is much more common than the latter, the system needs to show higher degree of robustness towards it and will therefore be penalized more if it fails to do so. A true detection, however, is more significant than either type of false detection and is therefore given more weight. The performance figures obtained for each of the datasets are listed in Table 4.

Dataset		System performance (S) in percent
PETS2006		85.00
PETS2007		75.00
AVSS i-LiDS	AB	83.33
	PV	37.50
	Overall	57.14
Custom	Outdoor	77.63
	Indoor	76.56
	Well Lit	80.76
	Dark	66.67
	Overall	72.72
All datasets combined		76.79

Table 4 Performance summary for different datasets.

As can be seen in Table 4, this system's performance was worst in the AVSS dataset, particularly in the parked vehicle (PV) dataset. The main reason for this was that the camera that recorded these videos was moving or shifting slightly every now and then, probably due to winds, thus leading to small changes in the viewpoint. Since this system tracks only static objects based on their precise position in the scene, such shifts caused the system to lose track of existing static

objects and start tracking them all over again. In addition, the camera exposure appeared to be changing frequently leading to several scene-wide lighting changes and this in turn caused the background to be reset repeatedly. Performance in the abandoned object (AB) dataset was much better with 100% detection rate though there were a couple of false detections in the AB-Hard video sequence.

The system's performance in the remaining datasets was fairly good, being consistently above the 70% accuracy mark. The only major failures in the custom dataset were in sets 7, 9, 15 and 18. The first one occurred because of extremely cluttered background that always presents problems for the edge detection based abandonment analysis employed here (Fig. 21). The second one occurred because the black colored test object was placed in a shadowy region and was invisible even to human eyes. A similar issue plagued the last two videos since they were shot in very dark indoor conditions and the test object simply blended in with the background in spite of applying contrast enhancement during pre processing.

Performance in the two PETS datasets was quite similar, with PETS2006 showing slightly better results. This occurred mainly because one of the views in PETS2007 (view 1) featured a very crowded queue that led to constant occlusions of the test object. In addition, the objects in this dataset were left behind for too short periods of times for alarm to be successfully raised.

5.5 Comparison with existing systems: Though both PETS2006 and AVSS i-LiDS datasets have been used for testing in many contemporary works in literature, very few of these report sufficiently detailed results for direct comparison with our results. For example, even though PETS2006 has been used in [9], [10], [21], [22], [24] and [29], the results for all 4 views have not been reported in even one. In addition, some, like [24] and [21], have provided no information about false positives.

Combined results for PETS2006 and AVSS-AB datasets have been reported in [9]. Using a similar but slightly more relaxed metric, the overall accuracy reported there is 85.20%. The corresponding figure for our system is 84.78%. though it must be noted that this has been obtained using a stricter metric. The results in [10] have been reported for AVSS-AB dataset with accuracy of 66.67% which is significantly worse than our result of 83.33%. The results of AVSS-PV and PETS2006 datasets have also been reported here but only partially and cannot be compared with our results. The system presented in [21] has been tested on PETS2006 and PETS2007 datasets but no information has been provided about false positives or individual views, again rendering it unsuitable for comparison with our system.

Results in [22], [24] and [29] have been reported for only one view for each of the 7 scenarios in PETS2006 dataset. In addition, the actual false positive count has not been reported in [22] and [24]. The approximate accuracy figures are 80% for [22] and 91.67% for [24] and 90% for [29]. These are comparable to our accuracy of 85% obtained over all the 4 views; choosing only the best result for each scenario would increase it to 100%. The system in [29] has also been tested

on AVSS dataset with accuracy of 58.93% which is only slightly better than our result of 57.14% in spite our system's failure to cope with the shaky nature of AVSS-PV dataset videos.




Fig. 21 Sample frame from set 7 of custom dataset. The current frame and background image are shown on the left (top and bottom) while the respective Canny edge detection results are on the right.

6. Conclusions and Future Work

The application was tested on a large number of publicly available as well as custom made videos and was found to give accuracy comparable to most contemporary systems, while still managing to run in real time on a fairly modest setup. One particularly noteworthy aspect of its performance was the very low rate of false positives that was obtained. Also, these results were obtained by using virtually the same methods and parameter values across all the videos, with only minor changes for some cases. Significantly better results can be obtained if exhaustive testing and adjustments are carried out for each scenario.

Finally, and perhaps most importantly, the methods that have been presented and tested in this paper are those that are part of the initial implementation of this system. These were specifically chosen to be relatively simple methods, both to implement and to execute, due to limitations of time and computational resources. However, owing to the modular nature of this system, it is quite easy to add more sophisticated methods to any of its module. This system can, therefore, be considered as the foundation for a truly robust framework that only requires a bit of calibration to perform well in practically any scenario.

Appendix A: Explanation of Source Code

The application has been coded entirely in C++ and makes use of OpenCV 2.4.4 library. The coding was done using Microsoft Visual Studio 2008 IDE and the application therefore exists as a Visual C++ solution.

A modular approach has been chosen to design this application with a view to facilitate easy expandability in the future. The application source is divided into 8 main modules and 4 helper modules, each consisting of one or more classes. A brief description of each of these modules follows along with an explanation of the important functions and variables.

A.1. Main modules

A.1.1 User interface module

Purpose: Creates the input-output interface for the user and allow the adjustment of various frame processing options (section 4.1).

Files: `userInterface.hpp` (header) and `userInterface.cpp` (function definitions)

Class Name: `UserInterface`

Core member functions:

- i. `initInputInterface`: sets up the input video stream either from camera or from a video file; also creates several windows with which the user interacts with the application.
- ii. `initOutputInterface`: sets up the video writers that save the outputs of various modules as a video file on disk; also sets up various images that display these outputs to the user in real time.
- iii. `initProcInterface`: sets up several images that are used for processing by other modules; also provides the user the option to change the processing resolution in real time.
- iv. `initInterface`: wrapper function for the above three functions and called whenever the user adjusts some input/processing parameters that require a system reset.
- v. `initIOParams/initProcParams`: initialize the input-output and processing parameters with values provided by the `SystemParameters` module.

Window related member functions:

- i. `initIOWindow/updateInputSource/updateIOParams`: setup the window that provides options for changing the input source in real time.
- ii. `initProcWindow/updateProcParams`: setup the window that allows the user to adjust the processing resolution in real time.

- iii. `selectObject`: creates a window that can be used to specify objects for other modules by drawing their bounding boxes using the mouse.
- iv. `getClickedPoint`: mouse callback function to capture the point clicked at by the user.
- v. `showIOHelpWindow/showProcHelpWindow`: display or remove the respective help windows.

A.1.2 Pre-processing module

Purpose: Provide functions to apply contrast enhancement and noise reduction to the input frames (sections 4.2 and 3.1).

Files: `preProcessing.hpp` (header) and `preProcessing.cpp` (function definitions)

Class Name: `PreProcessing`

Core member functions:

- i. `equalizeHistogramRGB`: perform histogram equalization (section 3.1.1.1).
- ii. `normalizeImageRGB/ normalizeImageGS`: perform contrast stretching for RGB and grayscale images respectively (section 3.1.1.2).
- iii. `filterImageRGB`: split rgb image into its constituent bands, apply filtering to each and combine the results (section 3.1.1.3).
- iv. `performContrastEnhancement`: wrapper function for the above methods.
- v. `performNoiseReduction`: apply smoothing filter to an image to reduce its white noise content (3.1.2)
- vi. `getImageHistogram`: compute and display the histogram for a grayscale image.
- vii. `initParams`: initialize the pre-processing parameters with values provided by the `SystemParameters` module.

Window related member functions:

- i. `initWindow/ updateParams/ updateContrastEnhancementMethod/ updateNoiseReductionMethod`: setup the window that provides the ability to toggle pre-processing steps and adjust the relevant parameters (section 4.2)
- ii. `selectObject`: creates a window that can be used to specify objects for other modules by drawing their bounding boxes using the mouse.
- iii. `getClickedPoint`: mouse callback function to capture the point clicked at by the user.
- iv. `showPreProcessHelpWindow`: display or remove the help window when user clicks on the main window.

A.1.3 BGS module

Purpose: Provide functions to conveniently use and switch between the various BGS methods adapted from the online OpenCV BGS library [50]. It therefore serves as a unifying frontend/wrapper for the different classes used for the four BGS techniques.

Files: `bgsModels.hpp` (header) and `bgsModels.cpp` (function definitions)

Class Name: `BgsModels`

Core member functions:

- i. `initBGSModels`: initializes the structures that hold parameters for each of the four BGS methods.
- ii. `initBGSMethod`: sets up the currently active BGS method.
- iii. `initBGSParams/initBGSToggleParams`: initializes the parameters of the chosen BGS method with values provided by the `SystemParameters` module.

Window related member functions:

- i. `initWindow/updateBGSMethod/updateBGSToggleParams`: setup the window that allows user to switch between BGS methods and adjust parameters of the chosen BGS method.
- ii. `showBGSHelpWindow`: display or remove the help window when user clicks on the main window.

A.1.4 Foreground analysis module

Purpose: Refine the foreground mask produced by the BGS module by removing shadows and regions of sudden lighting changes (section 3.3).

Files: `foregroundProcessing.hpp` (header) and `foregroundProcessing.cpp` (function definitions)

Class Name: `ForegroundProc`

Core member functions:

- i. `detectLightingChange`: detects and removes foreground regions created due to sudden lighting changes (section 3.3.1)
- ii. `detectShadowUsingComplexNCC`: detects shadows using the complex NCC approach (section 3.3.2.1)
- iii. `detectShadowUsingSimpleNCC`: detects shadows using the simple NCC approach (section 3.3.2.2)

- iv. `isShadow`: function pointer that can refer to either of the previous two functions.
- v. `removeFalseForeground`: wrapper function for the lighting change and shadow detection functions.
- vi. `isShadowCandidate`: decides if a pixel is a shadow candidate in simple NCC approach (section 3.3.2.2)
- vii. `detectFalseShadow`: refines the shadow candidates detected by the previous method (section 3.3.2.2).
- viii. `getWindowExtentForImage/getWindowExtentForPixel`: calculate the extents of the filtering window when it is centered at a particular pixel, over all pixels in the image.

Window related member functions:

- i. `initWindow/updateParams`: setup the window that allows the user to adjust the parameters of the currently active foreground analysis techniques (section 4.4).
- ii. `initMorphWindow/updateShadowMethod`: setup the window that allows the user to enable or disable the various foreground analysis techniques (section 4.4).
- iii. `showForegroundProcHelpWindow/showForegroundToggleHelpWindow`: display or remove the foreground processing parameters and toggles help windows when user clicks on the respective main window.

A.1.5 Blob extraction module

Purpose: Extract blobs (connected components) from the foreground mask (section 3.4). This module serves as a frontend for the blob extraction and manipulation library ‘cvBlobsLib [51].

Files: `blobDetection.hpp` (header) and `blobDetection.cpp` (function definitions)

Class Name: `ForegroundProc`

Core member functions:

- i. `getBlobs`: extracts blobs from the given binary image while discarding all that are below the minimum specified size; also displays the bounding boxes of these blobs overlaid on the current frame(section 3.3.1)
- ii. `initParams`: initializes the minimum blob size parameter.

Window related member functions:

- i. `initWindow/updateParams`: setup the window that can be used to adjust the minimum blob size (section 4.5).

- ii. `showBlobDetectionHelpWindow`: display or remove the help window for this module.

A.1.6 Blob tracking module

Purpose: Establish correspondence between new and existing blobs and use it to track static blobs in the scene (section 3.5).

Files: `blobTracking.hpp` (header) and `blobTracking.cpp` (function definitions)

Class Name: `BlobTracking`

Core member functions:

- i. `updateTrackedBlobs`: compares each of the blobs extracted from the current frame with the existing blobs to establish correspondence between them; also does occlusion detection (section 3.5.1)
- ii. `processTrackedBlobs`: updates the status/label of the tracked blobs on the basis of the result of the previous function; also adds and removes blobs from the tracking system (section 3.5.1-3.5.4).
- iii. `updateBlobImage`: update the image displaying the currently tracked blobs.
- iv. `updateTrackedStaticMask`: updates the mask showing all tracked objects visible in the current frame that is used as feedback for BGS module (section 3.5.5)
- v. `updateMovingAverageOfDifference`: updates a running average of pixel intensit difference for each tracked blob (section 3.5.3).
- vi. `initParamsMatch/ initParamsTrack`: initialize the matching and tracking parameters (section 4.6).

Window related member functions:

- i. `initWindowTrack/updateParamsTrack/updateStaticFactor`: setup the window that allows the user to adjust the parameters that control the duration for which blobs are to be tracked before changing their status or removing them from the system.
- ii. `initWindowMatch /updateParamsMatch/ updateDiffMethod/ updateDistanceMethod/ updateAreaMethod`: setup the window that allows the user to adjust the parameters that control the matching and correspondence of blobs (section 4.6).
- iii. `showTrackHelpWindow/ showMatchHelpWindow`: display or help window associated with the respective main window.

A.1.7 Abandonment analysis module

Purpose: Classify static blobs into either abandoned/removed objects, state changes or still persons (section 3.6)

Files: AbandonmentAnalysis.hpp (header) , abandonmentAnalysis.cpp (important function definitions), regionGrowing.cpp (helper function for region growing sub-module, mostly obsolete)

Class Name: ForegroundProc

Core member functions:

- i. detectRemovedObjectsUsingRegionGrowing: detects removed objects using the region growing technique (section 3.6.1.1).
- ii. detectRemovedObjectsUsingGradientEdgeDetection: detects removed objects using Sobel gradient based edge detection (section 3.6.1.2).
- iii. detectRemovedObjectsUsingCannyEdgeDetection: detects removed objects using Canny edge detection (section 3.6.1.2).
- iv. evaluateRegionGrowingPixelCount/performRegionGrowing: calculates the number of pixels added to a contour by applying region growing.
- v. evaluateGradientEdgePixelCount/evaluateCannyEdgePixelCount: calculates the number of pixels along an object's boundary that are edge pixels in the given image.
- vi. evaluatePixelSimilarity/comparePixelSimilarity/evaluateRegionSimilarity: implement the various metrics of evaluating a pixel's similarity to a region.
- vii. evaluateSimilarity : function pointer that can refer to any of the 3 previous functions.

Window related member functions:

- i. initWindow/updateParams/updateAbandonmentMethod/updateSimilarityMethod: setup the window that allows the user to adjust parameters pertaining to the currently active abandonment analysis method (section 4.7).
- ii. showAbandonmentHelpWindow: display or remove the abandonment analysis help window when user clicks on the main window.

A.1.8 Blob filtering module

Purpose: Allow the user to add objects and filter removed objects by comparing them to these added objects (section 3.7).

Files: blobFilter.hpp (header) and blobFilter.cpp (function definitions)

Class Name: BlobFilter

Core member functions:

- i. addCandidateRemovedObjects: create an interactive window where user can specify bounding boxes of objects to be added to the filtering system (section 4.8)
- ii. getObject: detect a single bounding box added by the user; also detect when existing objects are to be removed.
- iii. getClickedPoint: mouse callback function that detects mouse clicks and hover.
- iv. refreshSelectionImage: update the selection image to reflect the objects added so far.
- v. matchObjectLocation/ matchObjectSize/matchObjectAppearance: match two objects according to their location, size or appearance using user specified thresholds for each.
- vi. findMatchingObject: compare a given object with all the objects added to the filtering system to find a match.
- vii. filterRemovedObjects: uses the previous function to eliminate removed objects that do not match any existing filtering object.

Window related member functions:

- i. initWindow/updateParams/updateMatch/updateFiltering: setup the window that allows the user to adjust the blob filtering parameters (section 4.8).
- ii. initMorphWindow/updateShadowMethod: setup the window that allows the user to enable or disable the various foreground analysis techniques (section 4.4).
- iii. showBlobFilteringHelpWindow: display or remove the help window associated with this module.

A.2 Helper modules

A.2.1 System parameter initialization module

Purpose: Read the initial values of all the parameters for each of the main modules from a formatted ASCII file (called 'init_parameters.txt' by default) that makes it convenient for the user to specify these values.

Files: systemParameters.hpp (header) and systemParameters.cpp (function definitions)

Class Name: SystemParameters

Important member functions:

- i. `readInitialParams`: acts as a wrapper function for the individual parameter reading functions for each of the main modules.
- ii. `getParamString`: reads a single parameter as a string from the file while ignoring comments and blank spaces.
- iii. `getParamVal`: uses the previous function to obtain a numeric parameter value.

A.2.2 Combine images module

Purpose: Combine two images into a single composite image by stacking them horizontally or vertically; is used by the user interface module to combine several informative images into a single one for the user's convenience.

Files: `combineImages.hpp` (header) and `combineImages.cpp` (function definitions)

Class Name: There is no class in this module; it is merely a collection of related functions

Important functions:

- i. `initializeCombinedImage`: initializes an `IplImage` structure that can hold the result of combining the specified images in the specified configuration.
- ii. `combineImages`: combines two images into a single image.
- iii. `covertGrayscaleToRGB`: converts a grayscale image to an equivalent RGB image by using its single intensity value for all three channels. This is useful if grayscale and RGB images are to be combined by the last function since it requires all its input images to have the same number of channels.

A.2.3 Running statistics module

Purpose: Maintains the running mean and variance of a set of values whose members are added one at a time; it is used by the tracking module to maintain running average of differences (3.5.3)

Files: `runningStat.hpp` (header and definitions)

Class Name: There are two classes: `RunningStatVector` (for multi dimensional data) and `RunningStatScalar` (for single dimensional or scalar data)

Important functions:

- i. `addElement`: add an element to the set of values and simultaneously update their mean and variance.

A.2.4 BGS techniques modules

Purpose: These are four modules that implement the individual BGS techniques (section 3.2) and are all adapted from an online BGS library [50]. All of these classes are derived from a

single (virtual) base class called `Bgs` defined in `Bgs.hpp`. They also make use of some special data structures specified in the files `Image.hpp` and `Image.cpp`.

Files:

- Original GMM [3] (section 3.2.1): `GrimsonGMM.hpp` (header) and `GrimsonGMM.cpp` (function definitions)
- Improved GMM [4] (section 3.2.1): `ZivkovicAGMM.hpp` (header) and `ZivkovicAGMM.cpp` (function definitions)
- Adaptive median (section 3.2.2): `AdaptiveMedianBGS.hpp` (header) and `AdaptiveMedianBGS.cpp` (function definitions)
- Running Gaussian average (section 3.2.3): `WrenGA.hpp` (header) and `WrenGA.cpp` (function definitions)

Class Names:

- Original GMM: `GrimsonGMM`
- Improved GMM: `ZivkovicAGMM`
- Adaptive median: `AdaptiveMedianBGS`
- Running Gaussian average: `WrenGA`

Important functions: Following functions are common amongst all 4 classes:

- i. `Subtract`: this is the original background subtraction function that produces the foreground mask
- ii. `Subtract2`: this is a modified version of the previous one (used only in the two GMM methods) and uses a mask to selectively update background pixels.
- iii. `Update`: this is used in adaptive median and running Gaussian modules to selectively update the background model.
- iv. `pushIntoBackground`: pushes an object, specified by its bounding box, into the background by writing pixel values from the current frame into those locations.
- v. `resetBackground`: copies the current frame onto the background model except at the pixels specified through a mask.
- vi. `SubtractPixel/ SubtractPixel2`: original and modified versions of the pixel level background updating functions; the latter is defined only for the two GMM methods.

REFERENCES

- [1] N. McFarlane and C. Schofield, "Segmentation and tracking of piglets in images," *Machine Vision and Applications* Vol. 8, Issue 3, pp. 187–193, 1995.
- [2] C.Wren, A. Azarbayejani, T. Darrell, and A. Pentland, "Pfinder: real-time tracking of the human body," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, Vol. 19 Issue 7, pp. 780-785, July 1997.
- [3] C. Stauffer and W. E. L. Grimson, "Adaptive background mixture models for real-time tracking," in *Proc. IEEE Conference on Computer Vision and Pattern Recognition*, Vol. 2, pp. 246–252, Feb. 1999.
- [4] Z. Zivkovic, "Improved adaptive Gaussian mixture model for background subtraction," in *Proc. International Conference on Pattern Recognition*, Vol. 2, pp. 28–31, Aug. 2004.
- [5] A. Elgammal., D. Harwood and L.S. Davis, "Nonparametric background model for background subtraction," in *Proc. 6th European Conference on Computer Vision*, Vol. 2, pp. 751-767, 2000
- [6] A. Elgammal, D. Harwood, R. Duraiswami and L.S. Davis, "Background and Foreground Modeling Using Nonparametric Kernel Density Estimation for Visual Surveillance," in *Proc. The IEEE*, Vol. 90, Issue 7, pp.1151-1163, Jul. 2002.
- [7] N. M. Oliver, B. Rosario, and A. P. Pentland, "A Bayesian Computer Vision System for Modeling Human Interactions," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, Vol. 22, Issue. 8, pp. 831-843, Aug. 2000
- [8] B. Han, D. Comaniciu, and L.S. Davis, "Sequential kernel density approximation through mode propagation: applications to background modeling," in *Proc. Asian Conference on Computer Vision*, Jan. 2004.
- [9] A.Singh, S.Sawan, M.Hanmandlu, V.K.Madasu, B.C.Lovell "An Abandoned Object Detection System Based On Dual Background Segmentation," in *Proc. Sixth IEEE International Conference on Advanced Video and Signal Based Surveillance*, pp. 352-357, 2009.
- [10] Fatih Porikli, Yuri Ivanov and Tetsuji Haga "Robust Abandoned Object Detection Using Dual Foregrounds," *EURASIP Journal on Advances in Signal Processing*, 2008.
- [11] Y. Tian, M. Lu and A. Hapapur, "Robust and efficient foreground analysis for real-time video surveillance," in *Proc. IEEE Conference on Computer Vision and Pattern Recognition*, vol. 1, pp. 1182-1187, June 2005.
- [12] L. Li and M.K.H. Leung, "Integrating intensity and texture differences for robust change detection," *IEEE Transactions on Image Processing*, Vol. 11, Issue 2, pp. 105-112, Aug. 2002.

- [13] Q. Fan and S. Pankanti, "Robust Foreground and Abandonment Analysis for Large-Scale Abandoned Object Detection in Complex Surveillance Videos," in *Proc. Ninth IEEE International Conference on Advanced Video and Signal-Based Surveillance*, pp. 58-63, Sept. 2012.
- [14] J.C.S. Jacques, C.R. Jung and S.R. Musse, "Background Subtraction and Shadow Detection in Grayscale Video Sequences," in *Proc 18th Brazilian Symposium on Computer Graphics and Image Processing*, pp. 189-196, Oct. 2005.
- [15] J. Stander, R. Mech and J. Ostermann, "Detection of moving cast shadows for object segmentation," *IEEE Transactions on Multimedia*, Vol. 1, Issue 1, pp. 65-76, Mar. 1999.
- [16] N. Al-Najdawi, H. E. Bez, J. Singhai and E. A. Edirisinghe "A survey of cast shadow detection algorithms," *Pattern Recognition Letters*, Vol. 33, Issue 6, pp. 752-764, April 2012.
- [17] F. Chang, C. Chen and C. Lu, "A linear-time component labeling algorithm using contour tracing technique," *Computer Vision and Image Understanding*, Vol. 93, Issue 2, pp. 206-220, Feb. 2004.
- [18] A. Rosenfeld and P. Pfaltz, "Sequential operations in digital picture processing," *Journal of the Association for Computing Machinery*, Vol. 13, Issue 4, pp. 471-494, Oct. 1966.
- [19] C. Fiorio and J. Gustedt, "Two linear time Union-Find strategies for image processing," *Theoretical Computer Science*, Vol. 152, Issue 2, pp. 165-181, Feb. 1996.
- [20] R. Lumia, L. Shapiro and O. Zuniga, "A new connected components algorithm for virtual memory computers," *Computer Vision, Graphics, and Image Processing*, Vol. 22, Issue 2, pp. 287-300, May 1983.
- [21] Xuli Li, Chao Zhang, Duo Zhang "Abandoned Objects Detection Using Double Illumination Invariant Foreground Masks," in *Proc. 20th International Conference on Pattern Recognition*, pp. 436-439, Aug. 2010.
- [22] Smith, K., Quelhas, P. and Gatica-Perez, D, "Detecting Abandoned Luggage Items in a Public Space," in *Proc. Ninth IEEE International Workshop on Performance Evaluation of Tracking and Surveillance*, pp. 75 – 82, June 2006.
- [23] Z. Khan, T. Balch and F. Dellaert, "MCMC-based particle filtering for tracking a variable number of interacting targets," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, Vol. 27, Issue 11, pp. 1805-1819, Nov. 2005.
- [24] F. Lv, X. Song, X., B. Wu, V. Kumar and R. Singh, "Left Luggage Detection using Bayesian Inference," in *Proc 9th IEEE International Workshop on Performance Evaluation of Tracking and Surveillance*, pp. 83—90, June 2006.

- [25] R. Singh, S. Vishwakarma, A. Agrawal and M. D. Tiwari, "Unusual activity detection for video surveillance," in *Proc. 1st International Conference on Intelligent Interactive Technologies and Multimedia*, pp. 297-305, Dec. 2010.
- [26] D. W. Park, J. Kwon and K. M. Lee, "Robust visual tracking using autoregressive hidden Markov Model," in *Proc. IEEE Conference on Computer Vision and Pattern Recognition*, pp. 1964-1971, June 2012.
- [27] Z. Han, Q. Ye and J. Jiao, "Online feature evaluation for object tracking using Kalman Filter," in *Proc. 19th International Conference on Pattern Recognition*, pp. 1-4, Dec. 2008.
- [28] J. Connell, A.W. Senior, A. Hampapur, Y.-L. Tian, L. Brown and S. Pankanti, "Detection and Tracking in the IBM PeopleVision System," in *Proc. IEEE Conference on Multimedia and Expo*, Vol. 2, pp. 1403-1406, June 2004.
- [29] Y. Tian, R. Feris and A. Hampapur "Real-time detection of abandoned and removed objects in complex environments," in *Proc. IEEE International Workshop Visual Surveillance* (in conjunction with ECCV'08), Marseille, France, 2008.
- [30] Jianting Wen, Haifeng Gong, Xia Zhang and Wenzhe Hu "Generative model for abandoned object detection." *IEEE International Conference on Image Processing*, pp. 853-856, Nov. 2009.
- [31] Silvia Ferrando, Gianluca Gera, Massimo Massa, Carlo Regazzoni "A New Method for Real Time Abandoned Object Detection and Owner Tracking," *IEEE International Conference on Image Processing*, pp. 3329-3332, Oct. 2006.
- [32] Bhargava, M.; Chia-Chih Chen; Ryoo, M.S.; Aggarwal, J.K. "Detection of Abandoned Objects in Crowded Environments." In *Proc. IEEE Conference on Advanced Video and Signal Based Surveillance*, pp. 271-276, Sept. 2007.
- [33] Chi-Hung Chuang, Jun-Wei Hsieh, Luo-Wei Tsai, Sin-Yu Chen, and Kuo-Chin Fan "Carried Object Detection Using Ratio Histogram and its Application to Suspicious Event Analysis." *IEEE Transactions on Circuits and Systems for Video Technology*, pp. 911-916, June 2009.
- [34] Chuan-Yu Cho, Wei-Hao Tung and Jia-Shung Wang "A crowd-filter for detection of abandoned objects in crowded area." in *Proc. 3rd International Conference on Sensing Technology*, pp. 581-584, Nov. 30 – Dec. 3, 2008, Tainan, Taiwan.
- [35] Ying-li Tian, Rogerio Feris, Arun Hampapur "Robust Detection of Abandoned and Removed Objects in Complex Surveillance Videos." *IEEE Transactions on Systems, Man, and Cybernetics, Part C: Applications and Reviews*, Vol. 41, Issue 5, pp. 565-576, Sept. 2011.

- [36] Chih-Yang Lin and Wen-Hao Wang “An Abandoned Objects Management System Based on the Gaussian Mixture Model.” *International Conference on Convergence and Hybrid Information Technology*, 2008.
- [37] Simon Denman, Sridha Sridharan, Vinod Chandran “Abandoned object detection using multi-layer motion detection.” in *Proc. International Conference on Signal Processing and Communication Systems*, pp. 439-448, 2007.
- [38] E Auvinet, E Grossmann, C Rougier, M Dahmane, J Meunier “Left-luggage detection using homographies and simple heuristics.” in *PETS*, pp. 51–58, 2006.
- [39] Sadiye Guler, Jason A. Silverstein and Ian H. Pushee “Stationary objects in multiple object tracking.” in *Proc. IEEE Conference on Advanced Video and Signal Based Surveillance*, pp. 248-253, 2007.
- [40] M. Spengler and B. Schiele “Automatic detection and tracking of abandoned objects.” in *Proc. IEEE International Workshop VSPETS*, 2003.
- [41] Jiyan Pan, Quanfu Fan, Sharath Pankanti “Robust abandoned object detection using region-level analysis.” in *Proc. 18th IEEE International Conference on Image Processing*, pp. 3597-3600, Sept. 2011.
- [42] Xiaoyang Tan and Bill Triggs “Enhanced Local Texture Feature Sets for Face Recognition Under Difficult Lighting Conditions.” *IEEE Transactions on Image Processing*, Vol. 19 Issue 6, pp. 1635-1650, June 2010.
- [43] Kurt Skifstad, Ramesh Jain “Illumination independent change detection for real world image sequence.” *Computer Vision, Graphics, and Image Processing*, Vol. 46, Issue 3, pp. 387-399, June 1989.
- [44] Ying-li Tian , Rogerio Feris , Arun Hampapur “Modeling of temporarily static objects for robust abandoned object detection in urban surveillance.” in *Proc. Eight IEEE International Conference on Advanced Video and Signal-Based Surveillance*, pp. 36-41, Sept. 2011.
- [45] N. Bird, S. Atev, N. Caramelli, R. Martin, O. Masoud and N. Papanikolopoulos, “Real Time, Online Detection of Abandoned Objects in Public Areas.” in *Proc. IEEE International Conference on Robotics and Automation*, pp. 3775-3780, May 2006.
- [46] J. Canny, “A Computational Approach to Edge Detection,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, Vol. 8, Issue 6, pp. 679-698, Nov. 1986.
- [47] “PETS 2006 Benchmark Data,” <http://www.cvg.rdg.ac.uk/PETS2006/data.html>, June 2013.
- [48] “PETS 2007 Benchmark Data,” <http://www.cvg.rdg.ac.uk/PETS2007/data.html>, June 2013.

[49] "2007 IEEE International Conference on Advanced Video and Signal based Surveillance (AVSS 2007)," http://www.eecs.qmul.ac.uk/~andrea/avss2007_d.html, June 2013.

[50] OpenCV based BGS Library written by Donovan Parks and available at <http://dparks.wikidot.com/source-code>.

[51] OpenCV based blob extraction library "cvBlobsLib" available at <http://opencv.willowgarage.com/wiki/cvBlobsLib>.